

Storage Systems (StoSys)

XM_0092

Lecture 5: Key-Value Stores

Animesh Trivedi
Autumn 2020, Period 2



Reminder: for the coming weeks

We will be gradually transforming to networking and distributed systems

It is important you understand networking basics and important concepts such as

- TSO, LRO, Jumbo Frames, Multicore scalability, affinities, and RDMA, etc.

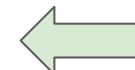
I will only introduce these topics selectively

Background reading: Please check out lecture 1, 2 (networking basic), 4 (multicore scalability), and 6 (RDMA networking) from the networking course linked below

- Slides are cross-uploaded in the Canvas for the Storage course
- The course page, Advanced Network Programming

Syllabus outline

1. ~~Welcome and introduction to NVM (today)~~
2. ~~Host interfacing and software implications~~
3. ~~Flash Translation Layer (FTL) and Garbage Collection (GC)~~
4. ~~NVM Block Storage File systems~~
5. NVM Block Storage Key-Value Stores
6. Emerging Byte-addressable Storage
7. Networked NVM Storage
8. Trends: Specialization and Programmability
9. Distributed Storage / Systems - I
10. Distributed Storage / Systems - II

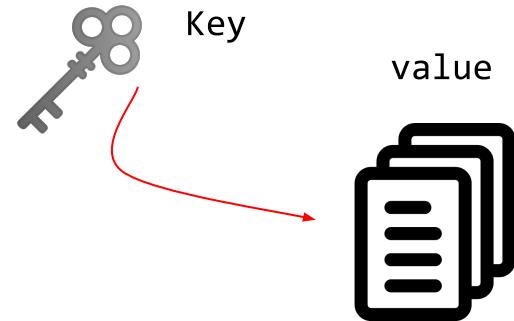


So, What is a Key-Value Store

A simplified data structure to store data and identify with a key

Examples: associate arrays (array?), dictionaries, hash table

Quite popular with web, scalable services



Isn't a file system suppose to store our data?

- FSes create new files, directories for every object
- Web objects are often small, but basic file system inode overheads per directory/files
 - inodes can be a few kB, if you want to store 64 bytes of data?
- Files/directories are difficult to iterate over quickly
- Range based queries need further indexing
- Object stores can support flexible consistent models (with FSes typically is a bad idea)
- Performance and feature optimizations, e.g., deduplication, transactions, compression, etc.

Basic Operations

`put(key, value)` : saves a value associated with a key

`value = get (key)` : retrieve the value associated with a key

`delete(key)` : deletes a key (can be equivalent of `put(key, NULL)`)

Batch versions of these commands: `multiget`, `multiput`

Range based queries: `iterate (start_key, end_key);`

Further helper commands: `replace`, `add`, `incr`, `decr`, etc.

No single data structure can do all operations efficiently

Layout of Coming Slides

B+ Trees and what they are good for

- What you need to do for storing them efficiently on NAND flash

LSM tree based KV design

- The basic idea
- LSM trees on OCSSD
- Application amplification in LSM

[Optional] A **Hash table**-based KV design (see Backup slides)

- FlashStore (and general topic of {memory \longleftrightarrow I/O} tradeoff)

B+ Tree

M-ary tree with sorted (keys-values) stored in leaves

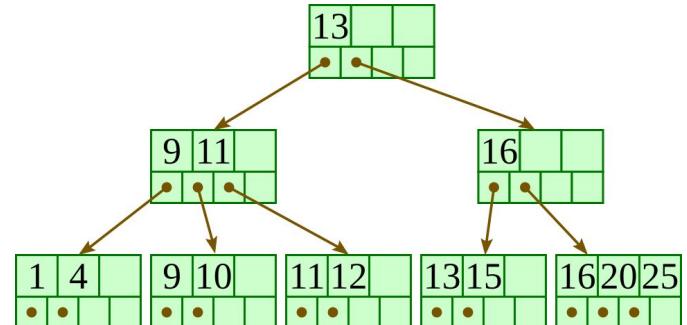
Useful for block-storage devices as it facilitate on-demand node fetching from the storage in block granularity

d-order tree has “d” keys and (d+1) pointers in non-leaf nodes, non-leaf nodes only contains “keys” for pivoting

Self-balancing (by splitting and merging nodes) and distance to all leaves nodes are equal from the root

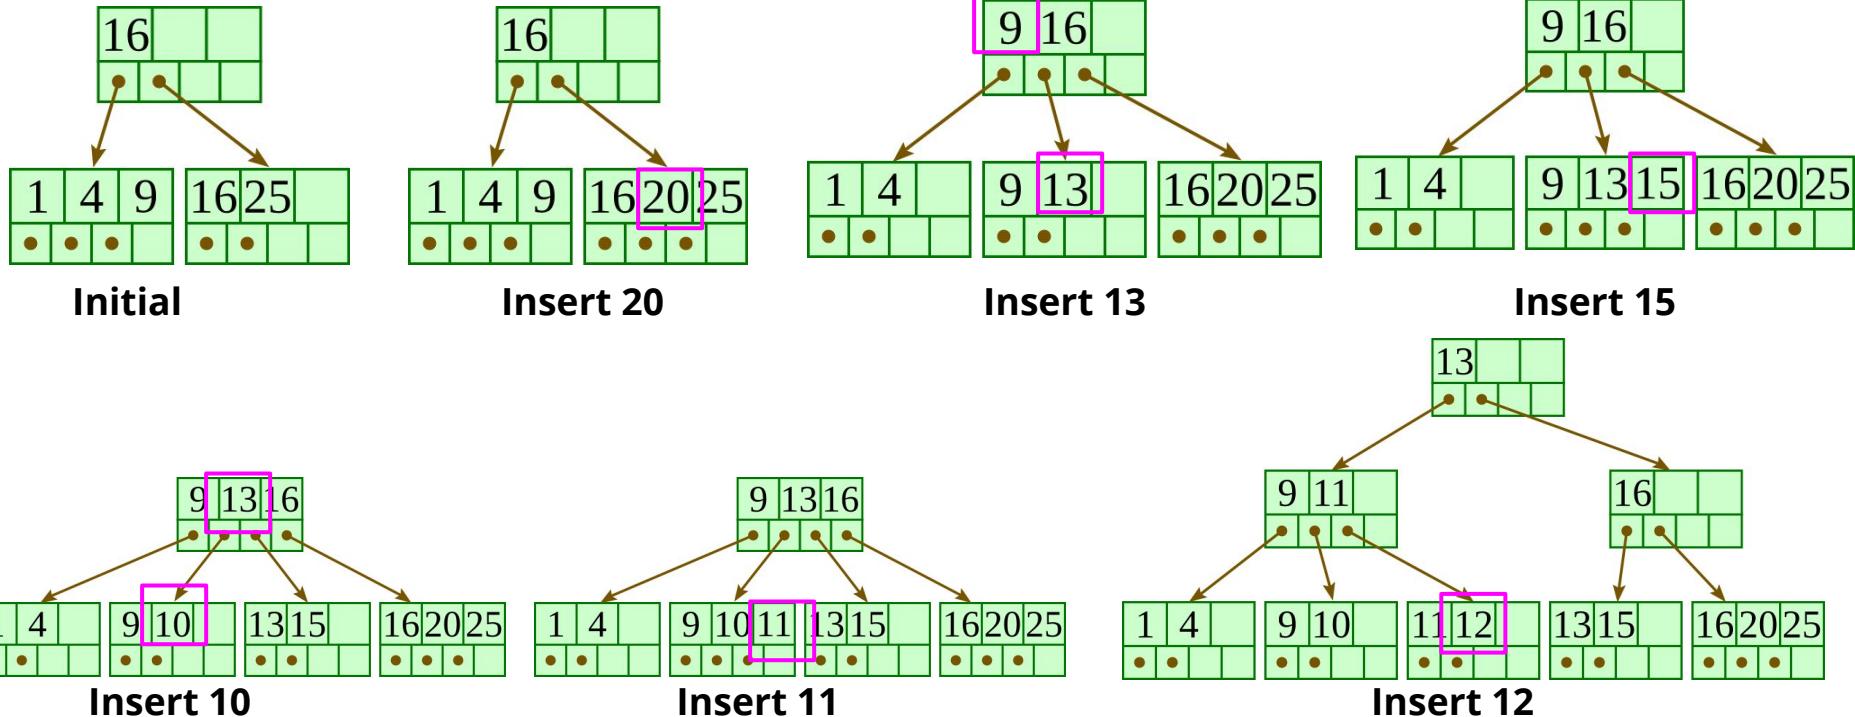
Popular data structure, used in Databases (Oracle, SQL) and file systems (ext4)

Optimized for read-heavy workloads (sorted indexes)

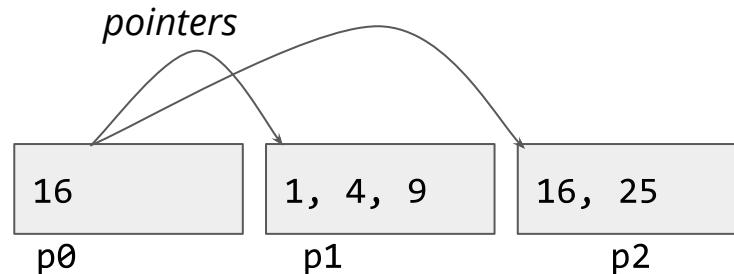
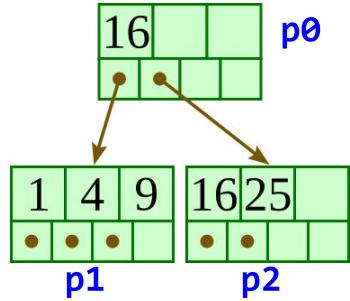


<http://www.cchurch.com/cs/340/reading/btree/index.html>

Example: B+ Tree Insertions



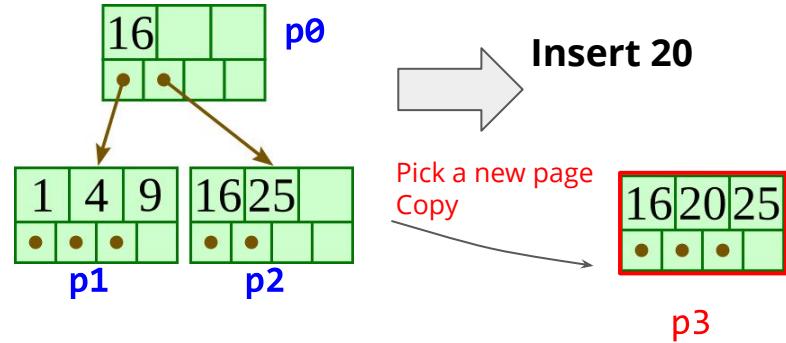
Example: B+ Tree Insertions on NAND Flash



NAND flash pages, the same layout used with HDD too

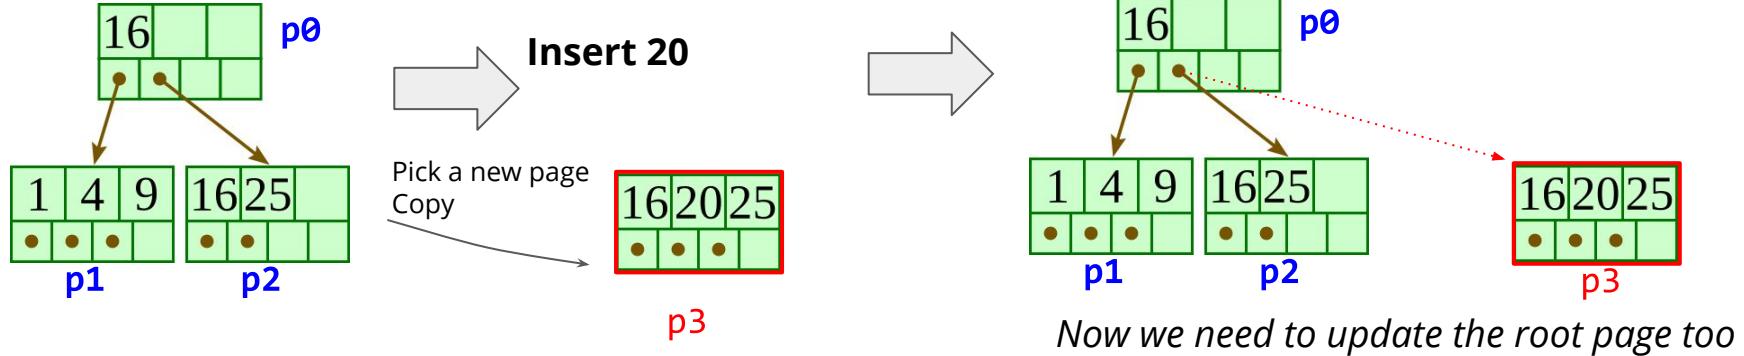
- Whole pages can be read in a single go
- Large sequential transfers, good performance
- All values sorted, so we know which page to load for which node

Example: B+ Tree Insertions on NAND Flash

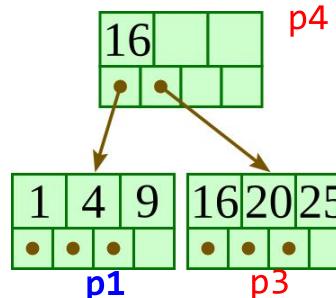
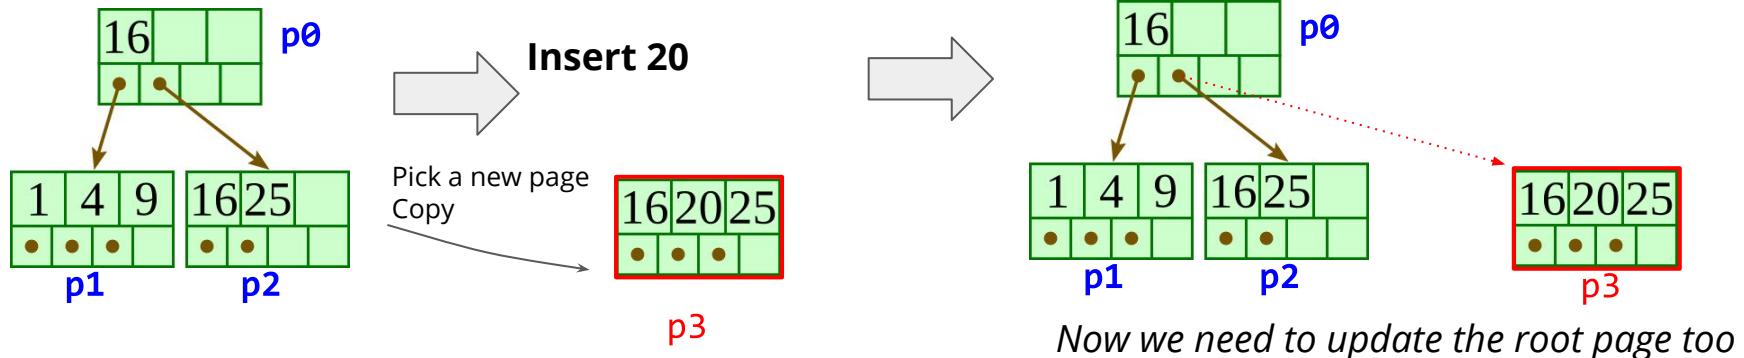


NAND pages cannot be in-place updated

Example: B+ Tree Insertions on NAND Flash



Example: B+ Tree Insertions on NAND Flash



For a simple value insertion we ended up writing 2 new pages (p_3 and p_4) and generating 2 old (p_0 and p_2), invalid pages

In general, for a tree " H " height, we will have to update all pages on the path to the root, " H " nodes

It's the same problem what we saw in Log-Structured FS (recursive update problem or also known as Wandering Tree problem)

B+ Trees on NAND Flash

μ -Tree : An Ordered Index Structure for NAND Flash Memory

Dongwon Kang

Dawoon Jung

Jeong-Uk Kang

Jin-Soo Kim

Computer Science Division

Korea Advanced Institute of Science and Technology (KAIST)

Daejeon 305-701, Korea

{dwkang,dwjung,ux}@camars.kaist.ac.kr

jinsoo@cs.kaist.ac.kr

ABSTRACT

As NAND flash memory becomes increasingly popular as data storage for embedded systems, many file systems and database management systems are being built on it. They require an efficient index structure to locate a particular item quickly from a huge amount of directory entries or database records. This paper proposes μ -Tree, a new ordered index structure tailored to the characteristics of NAND flash memory. μ -Tree is a balanced tree similar to B⁺-Tree. In μ -Tree, however, all the nodes along the path from the root to the leaf are put together into a single flash memory page in order to minimize the number of flash write operations when a leaf node is updated. Our experimental evaluation shows that μ -Tree outperforms B⁺-Tree by up to 28% for traces extracted from real workloads. With a small in-memory cache of 8 Kbytes, μ -Tree improves the overall performance by up to 90% compared to B⁺-Tree with the same cache size.

Categories and Subject Descriptors

H.3.1 [Content Analysis and Indexing]: Indexing methods; D.4.3 [File Systems Management]: Directory structures

General Terms

Algorithms, Design, Performance

Keywords

B⁺-Tree, NAND Flash, index structure

1. INTRODUCTION

Flash memory is being widely adopted as a storage medium for many portable embedded devices such as PMPs (portable media players), PDAs (personal digital assistants), digital cameras and camcorders, and cellular phones. This is mainly due to the inherent advantageous features of flash memory: non-volatility, small and lightweight form factor, low-power consumption, and solid state reliability.

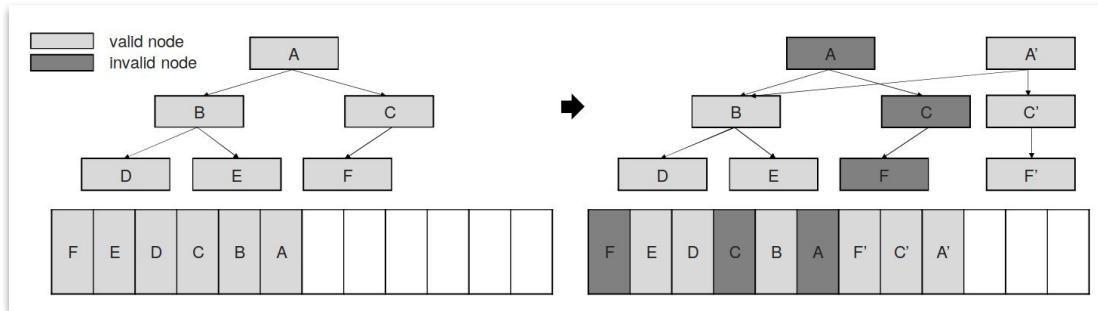
Flash memory comes in two flavors. The NOR type is usually used for storing codes since it can be directly addressable by processors. On the other hand, the NAND type is accessed on a page basis (typically 512 bytes ~ 4 Kbytes) and provides higher cell densities. The NAND type is primarily used for removable flash cards, USB thumb drives, and internal data storage in portable devices.

As the NAND flash technology development continues to double density growth on an average of every 12 months [23], the capacity of a single NAND chip is getting larger at an increasingly lower cost. The declining cost of NAND flash memory has made it a viable and economically attractive alternative to hard disk drives especially in portable embedded systems. As a result, many flash-aware file systems and embedded database management systems (DBMSs) are currently being built on NAND flash memory [2, 7, 9, 13, 24].

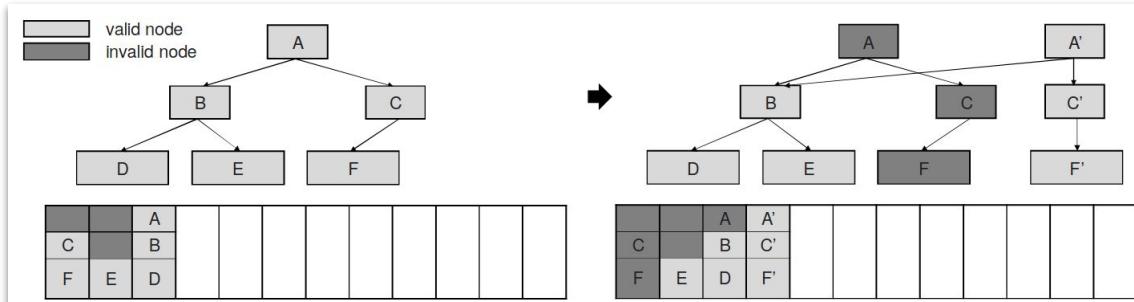
Any file system or DBMS requires an efficient index structure to locate a particular item quickly from a huge amount of directory entries or database records. For small scale systems, the index information can be kept in main memory. For example, JFFS2 keeps the whole index structures in memory that are necessary to find the latest file data on flash memory [24]. Apparently, this approach is not scal-

μ -Tree : The Basic Idea

Rearrange the layout, do not give each nodes its own page. Store multiple nodes on a single page: typically along the path which will be updated in case of an insertion



Basic ("N" writes)



Proposed (update in 1 write)

How to Pack Nodes in a Page

Should we equally divide space in a page to all levels

Keeps the logic simple, and searchable, we will know exactly which offset in a page a level starts

However,

- Then we need to “fix” the maximum height of the tree
- Key space exponentially increases at every level
 - L3 : 2 order tree with 3 pointers
 - L2 : 3×3
 - L1 : $3 \times 3 \times 3$ pointers ← this contain data, so we need to proportionally distribute space for different levels with flexibility to increase the level as we increase (or decrease the size of the tree)



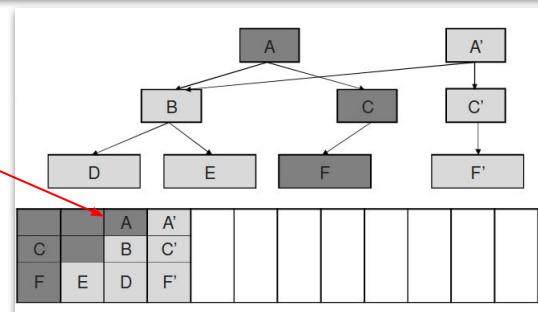
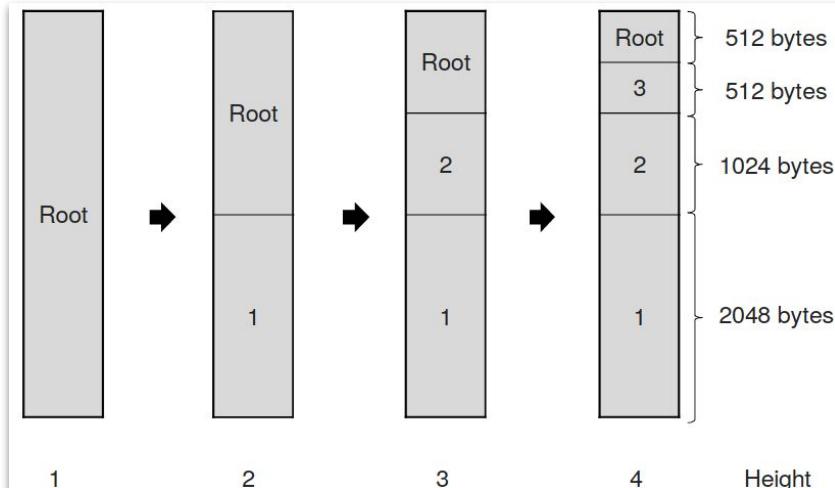
μ -Tree: Proportional Packing

In this setup

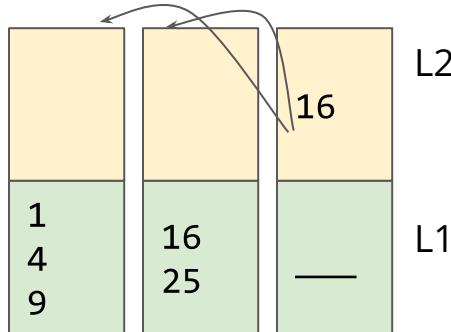
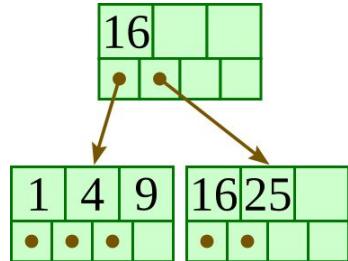
- Nodes within a page are still searchable
 - For a given level, and the height of the tree I can calculate which offset the node data starts
- Proportionally distribute space to different levels
- Enables us to do updates in one go, while keeping some data in old pages

The only thing we need to keep track of which page contains the “Root” pointer

- Changed from p2 to p3

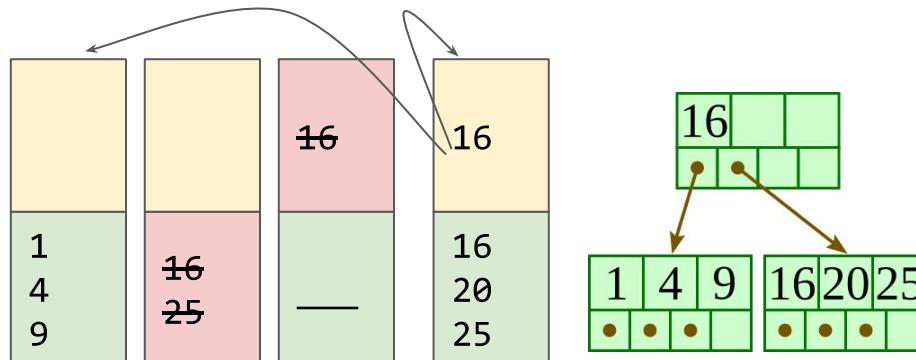


μ -Tree Insertions on NAND Flash



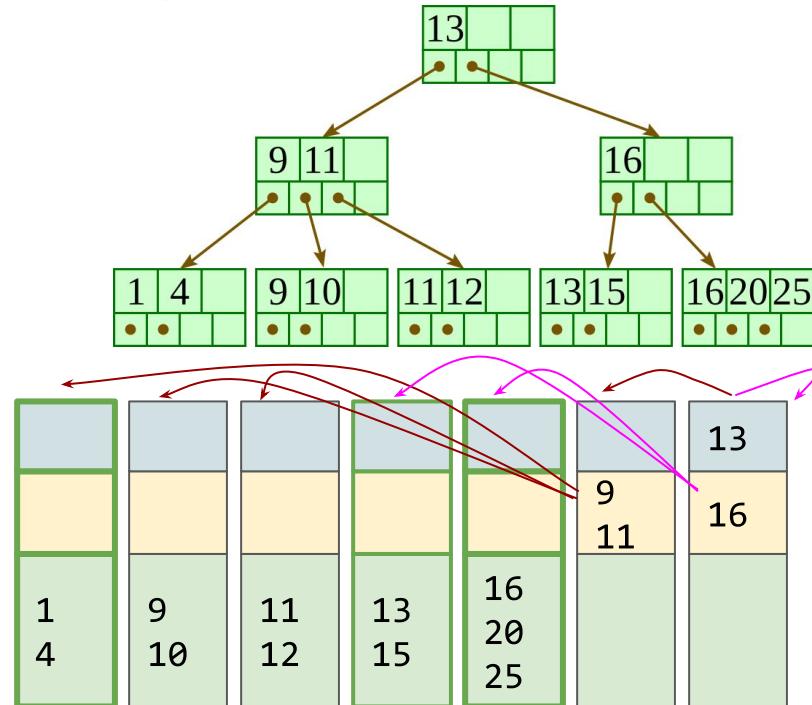
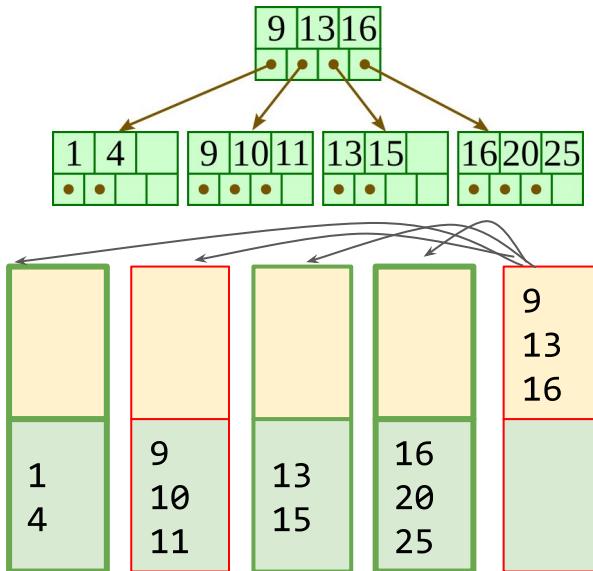
- In this case:
- 2 pages reading
 - 1 page writing

Insert 20



In general: $H \times \text{reading} + 1 \times \text{writing}$

μ -Tree Insertions with Height Increase



Eventually as you write more, things will be grouped together (the update path) on the same page blocks. A similar logic applies to deletion and tree compaction logic (skipped).

μ -Tree: Performance (analytical)

Since the number of pointers that can be stored in a single page for a given level is different for μ and B+ Trees

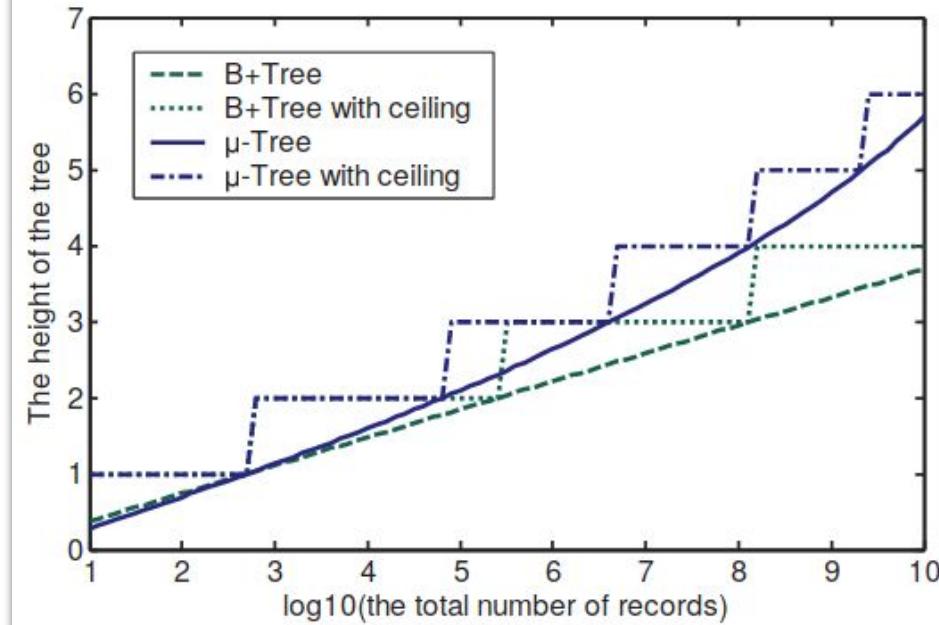
- Height difference, within +1 (upto 1B)
- Takes twice as much flash space

Will result in more reads

Table 3: The cost of operations

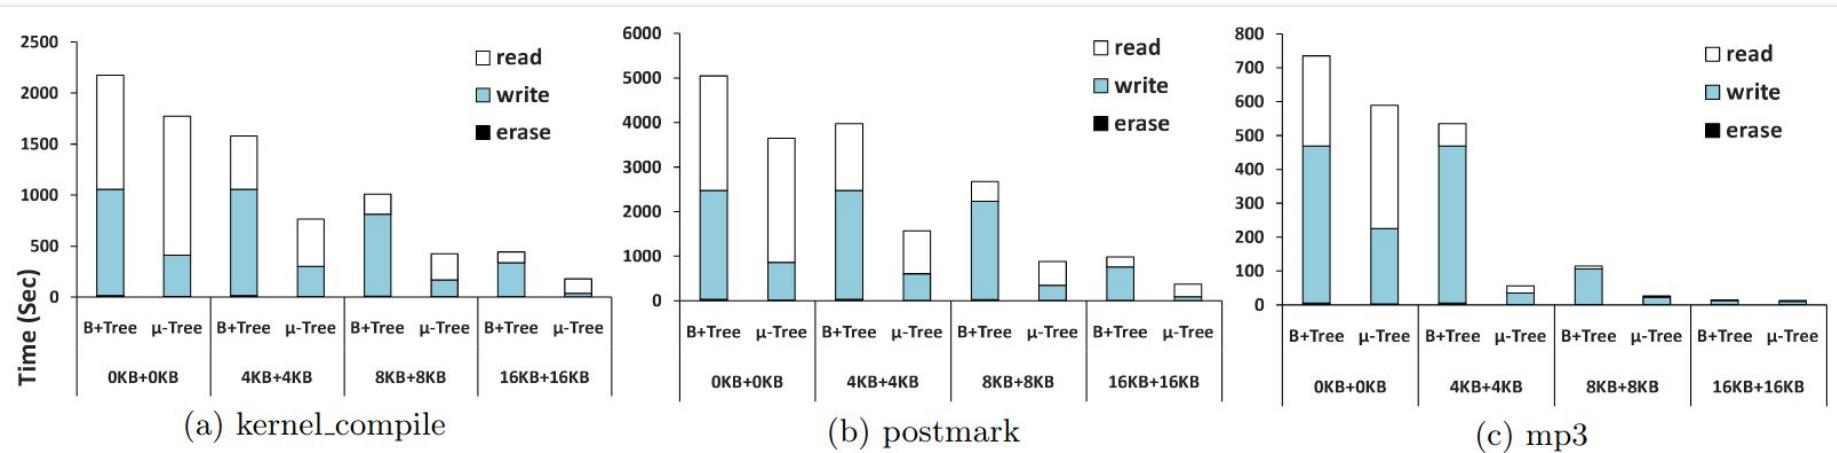
Operations	B ⁺ -Tree	μ -Tree
Retrieval	$c_r h_B$	$c_r h_\mu$
Insertion	$(c_r + c_w)h_B$	$c_r h_\mu + c_w$
Deletion	$(c_r + c_w)h_B$	$c_r h_\mu + c_w$

In absence of a split or collapse



μ -Tree: Performance

Traces collected from ReiserFS (B+ tree) about node creation, access, deletions
Could have used some other benchmarks (well!)



Better performance : decreases the number of writes and with more reads (taller tree)

There are other works too

An Efficient B-Tree Layer Implementation for Flash-Memory Storage Systems

CHIN-HSIEN WU and TEI-WEI KUO

National Taiwan University

and

LI PING CHANG

National Chiao-Tung University

With the significant growth of the markets for consumer electronics and various embedded systems, flash memory is now an economic solution for storage systems design. Because index structures require intensively fine-grained updates/modifications, block-oriented access over flash memory could introduce a significant number of redundant writes. This might not only severely degrade the overall performance, but also damage the reliability of flash memory. In this paper, we propose a very different approach, which can efficiently handle fine-grained updates/modifications caused by B-tree index access over flash memory. The implementation is done directly over the flash translation layer (FTL); hence, no modifications to existing application systems are needed. We demonstrate that when index structures are adopted over flash memory, the proposed methodology can significantly improve the system performance and, at the same time, reduce both the overhead of flash-memory management and the energy dissipation. The average response time of record insertions and deletions was also significantly reduced.

Categories and Subject Descriptors: C.3 [**Special-Purpose and Application-Based Systems**]: Real-Time and Embedded Systems; H.3.1 [**Content Analysis and Indexing**]: Indexing Methods; H.3.3 [**Information Search and Retrieval**]: Search Process

General Terms: Design, Performance, Algorithm

Additional Key Words and Phrases: Flash memory, B-tree, storage systems, embedded systems, database systems

ACM Reference Format:

Wu, C.-H., Kuo, T.-W., and Chang, L.-P. 2007. An efficient B-tree layer implementation for flash-memory storage systems. ACM Trans. Embedd. Comput. Syst. 6, 3, Article 19 (July 2007), 23 pages. DOI = 10.1145/1275986.1275991 <http://doi.acm.org/10.1145/1275986.1275991>

FlashDB: Dynamic Self-tuning Database for NAND Flash

Suman Nath
Microsoft Research
sumann@microsoft.com

Aman Kansal
Microsoft Research
kansal@microsoft.com

ABSTRACT

FlashDB is a self-tuning database optimized for sensor networks using NAND flash storage. In practical systems flash is used in different packages such as on-board flash chips, compact flash cards, secure digital cards and related formats. Our experiments reveal non-trivial differences in their access costs. Furthermore, databases may be subject to different types of workloads. We show that existing databases for flash are not optimized for *all* types of flash devices or for *all* workloads and their performance is thus suboptimal in many practical systems. FlashDB uses a novel self-tuning index that dynamically adapts its storage structure to workload and underlying storage device. We formalize the self-tuning nature of an index as a two-state task system and propose a 3-competitive online algorithm that achieves the theoretical optimum. We also provide a framework to determine the optimal size of an index node that minimizes energy and latency for a given device. Finally, we propose optimizations to further improve the performance of our index. We prototype and compare different indexing schemes on multiple flash devices and workloads, and show that our indexing scheme outperforms existing schemes under *all* workloads and flash devices we consider.

Categories and Subject Descriptors: H.2.4 [**Database Management Systems**]: Query processing; H.3.1 [**Content Analysis and Indexing**]: Indexing methods

General Terms: Algorithms, Design, Measurement, Performance.

Keywords: B⁺-tree, NAND Flash, indexing, log-structured index,

example includes sensor networks of mobile devices which have significant local processing power [4, 12]. In these cases rather than uploading the entire raw data stream, one may save energy and bandwidth by processing queries locally at a cluster-head or a more capable node and uploading only the query response or the compressed or summary data. Storage centric networks have also been discussed in [6, 7].

In most cases where the storage is part of the sensor network, the storage device used is flash based rather than a hard disk due to shock resistance, node size, and energy considerations. Additionally, flash is also common in many mobile devices such as PDA's, cell-phones, music players, and personal exercise monitors. These devices can benefit from a having light weight database.

Our objective is to design storage and retrieval functionality for flash storage. A simple method is to archive data without an index, and that is in fact efficient in many scenarios. However, as we show in section 6, for scenarios where the number of queries is more than a small fraction ($\approx 1\%$) of the number of data items, having an index is useful. Hence, we focus on indexed storage. Prior work on flash storage provides file systems (e.g., ELF [5]) and other useful data structures such as stacks, queues and limited indexes (e.g., Capsule [14], MicroHash [22]). Our goal is to extend the functionality provided by those methods to B⁺-tree based indexing to support useful queries such as lookups, range-queries, multi-dimensional range-queries, and joins.

Existing database products are not well suited for sensor networks due to several reasons. Firstly, existing products, including

Now, what about write-heavy workloads?

Write heavy workloads on flash can be really bad

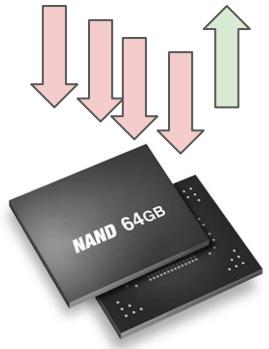
- Key-Values can be really small (32-64-128 bytes). Hash checksums are also small (512/1024 bits)

The best solution so far we have seen is a log (FTL, file system)

- Append small writes to a log and read from there (search)

How can we improve searching the log?

- We can build a hash table (key) → {flash offset}
 - But will need a lot of memory for the hash table
 - Simply 8 bytes per key (similar to FTL)
- Does not allow doing fast range-based queries and lookups **ideas?**



Back to the Future: LSM Trees

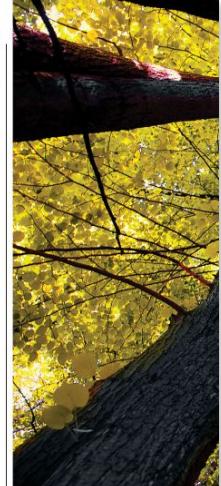
Log-Structured Merge (LSM) Tree data structure

Invented and optimized for HDD, why?

- Same logic as LogFS
 - Disks have fast sequential performance
 - Disks have poor random, small I/O performance
- Read/Write large chunks to disk
- Eliminates random insertions, updates and deletions

*Patrick E. O'Neil, Edward Cheng, Dieter Gawlick, Elizabeth J. O'Neil:
The Log-Structured Merge-Tree (LSM-Tree). Acta Informatica 33(4):
351-385 (1996)*

Very popular data structure: Bigtable, HBase, LevelDB,
SQLite4, Tarantool, RocksDB, WiredTiger, Apache Cassandra,
InfluxDB, and ScyllaDB



Article development led by  queue.acm.org

DOI:10.1145/3209210

Different uses for read-optimized B-trees and write-optimized LSM-trees.

BY ALEX PETROV

Algorithms Behind Modern Storage Systems

when they arise, and fine-tunes the database for your workload.

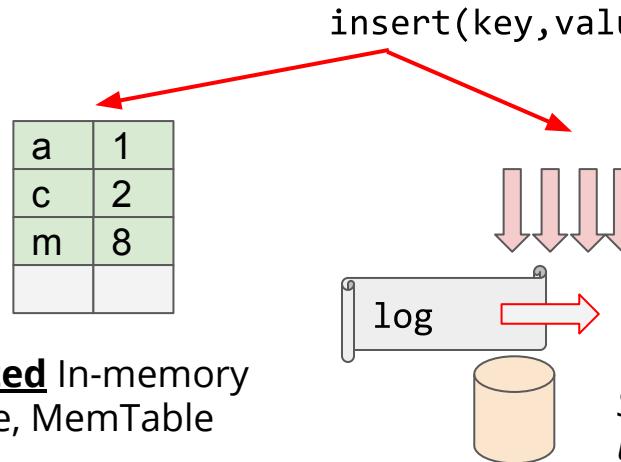
It is impossible to optimize a system in all directions. In an ideal world there would be data structures guaranteeing the best read and write performance with no storage overhead but, of course, in practice that is not possible.

This article takes a closer look at two storage system design approaches used in the majority of modern databases—read-optimized B-trees and write-optimized LSM (log-structured merge-trees)—and describes their use cases and trade-offs.

D. T. T. T.

<https://queue.acm.org/detail.cfm?id=3220266>

LSM Tree Basics



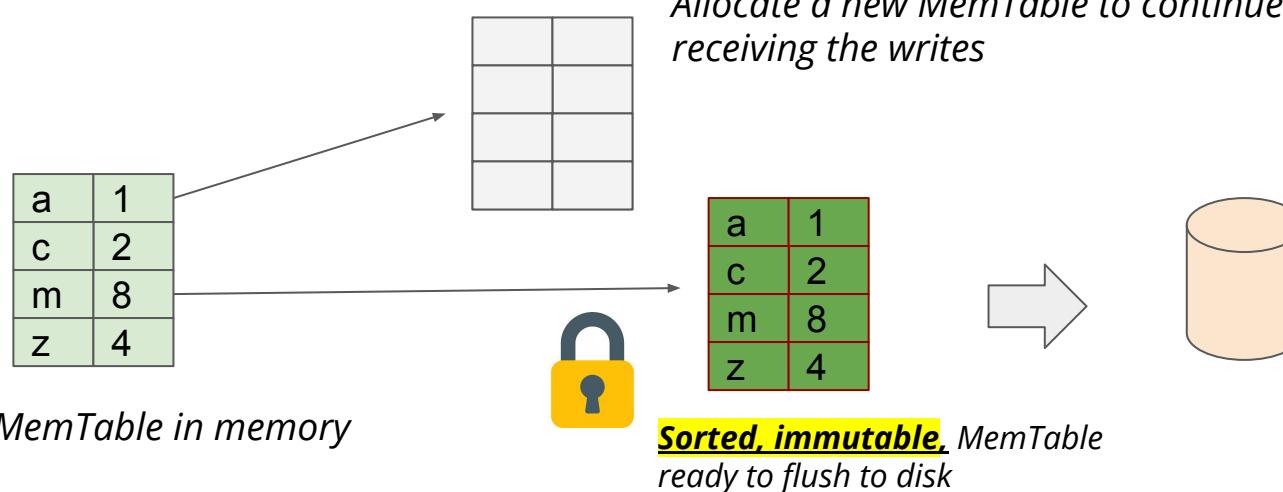
At insertion, (key,value) is

- written to the disk-resident log (large sequential performance)
- Inserted in the sorted MemTable to enable fast lookup with a range based query

Sequential log on disk, only used in for failure recovery

What happens when the in-memory data structure is full?

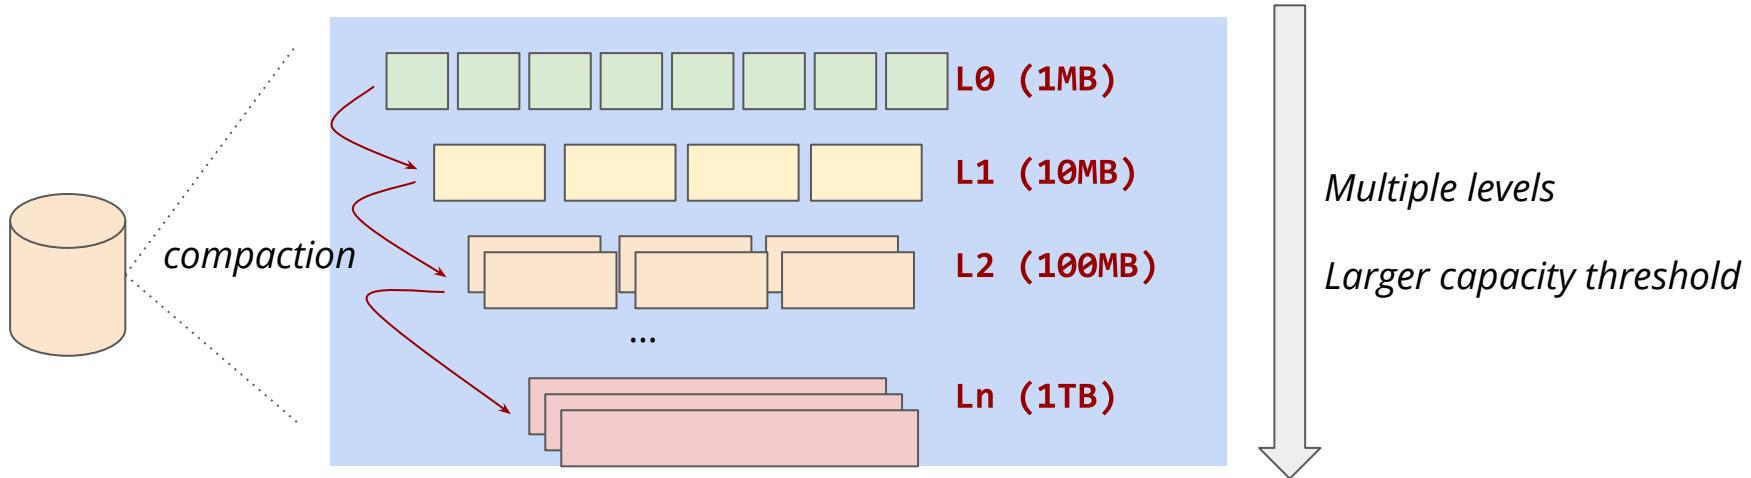
LSM Tree Basics



- Once the in-memory table is full : the MemTable is marked immutable and *flushed* to disk
- Key lookup requires searching in the MemTable + looking up on the disk
 - (we will see how this can be made efficient)
- If data is present in both location, use timestamps to reconcile which is the newest write

Challenge now is how to (a) manage and (b) search TBs of data on disk to look for a key

LSM Tree Basics



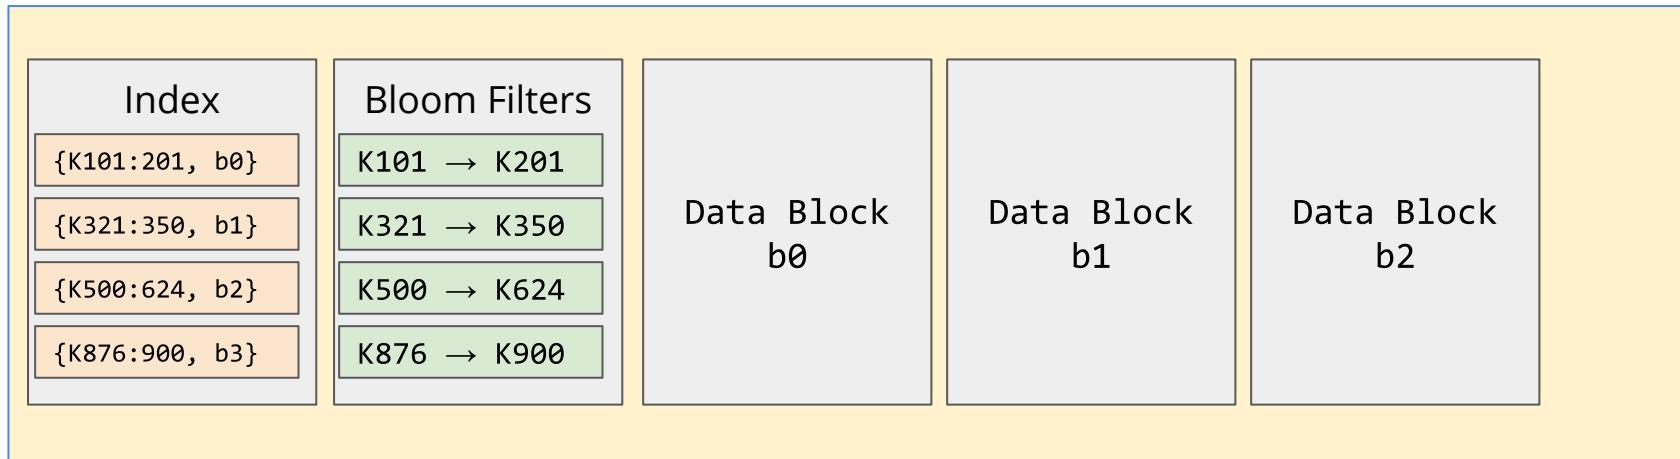
Data is stored in a multi-level, large, immutable files on the disk (no holes/gaps). Each level has a fixed size that increases as you go to the higher levels

A new table flush is written always written to **L0**

Just like in-memory table, once, a preconfigured size of file is reached, files at level i can be merged with $(i+1)$. This process is known as **compaction**. Since files written are sorted, the compaction is essentially an N-way merge sort from level (i) to $(i+1)$

On-Disk File Format (SSTables)

Sorted String Tables (SSTables)



When searching : find a value in the index range, then check in the bloom filter

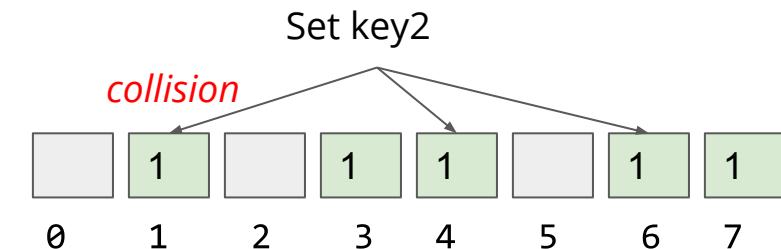
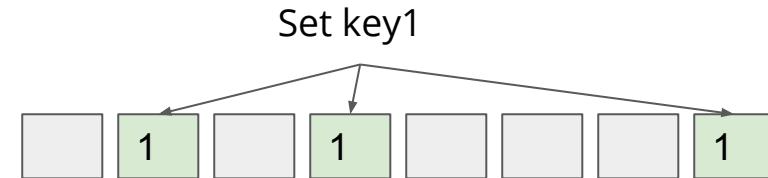
Then go fetch the “block” for reading and scan the value inside

All files are immutables, hence, a delete is a new insertion with a “NULL” value

Recap: Bloom Filters



A bunch of hash function, **h1, h2, h3**



Now if we were to check for key3 and key4

- lookup (key3) = 1, 4, 7 // all set, but they key3 was never set, **false positive**
- lookup (key4) = 0, 2, 5 // nope, this key was never set, always accurate! **Cannot have false negative**

The rate of false positive depends upon the size of the filter and the quality of the hash functions

Example Compaction Process

L0

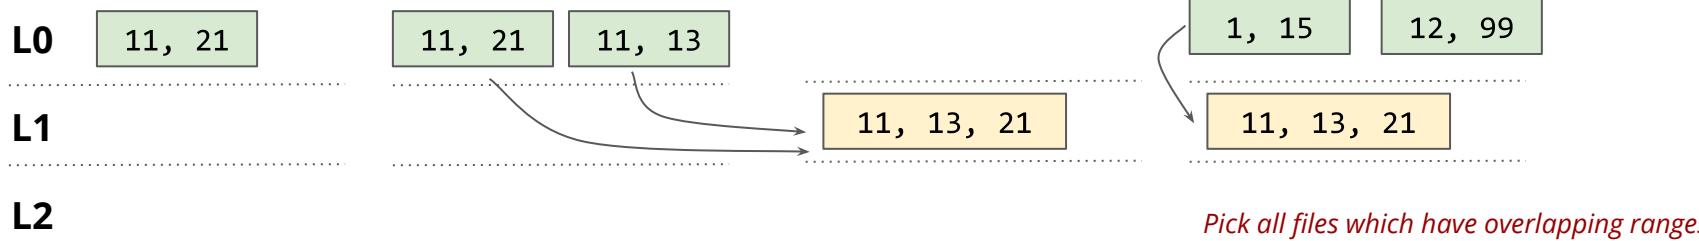
11, 21

.....
L1

.....
L2

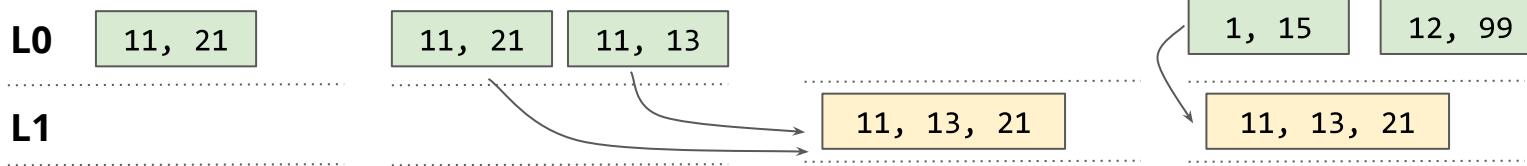
Example Compaction Process

L0 can have duplicates keys in different files

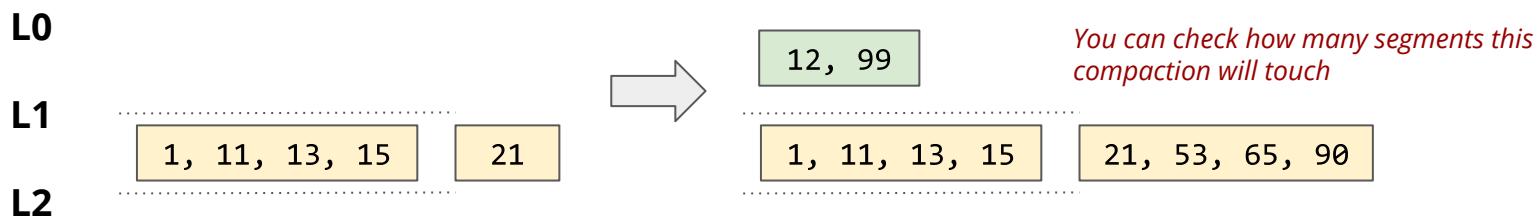


Example Compaction Process

L0 can have duplicates keys in different files

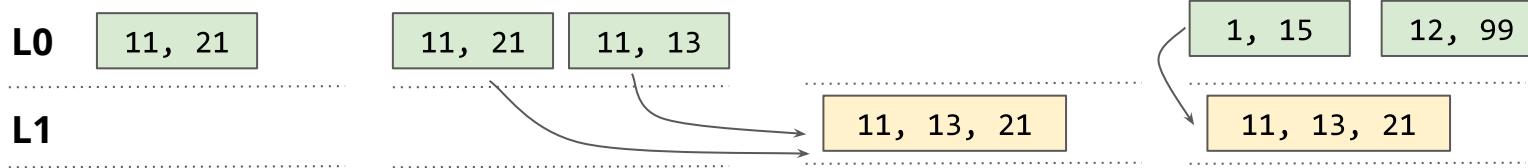


Pick all files which have overlapping ranges

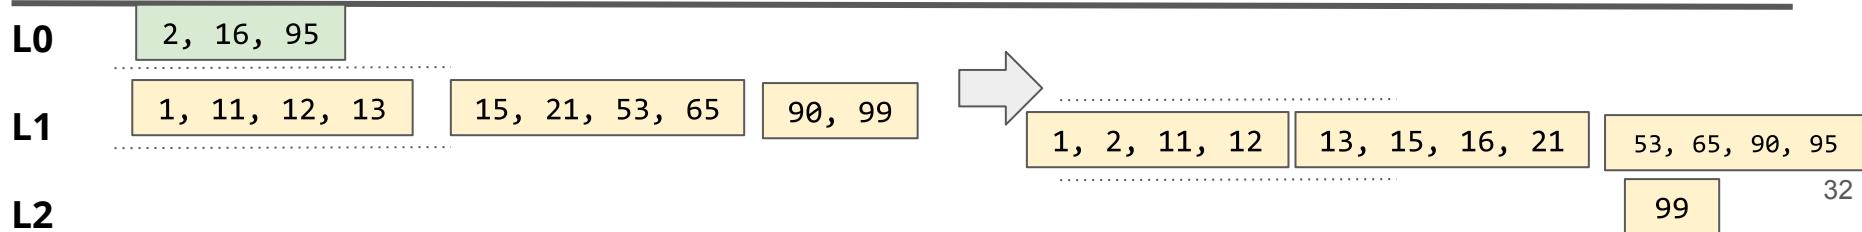
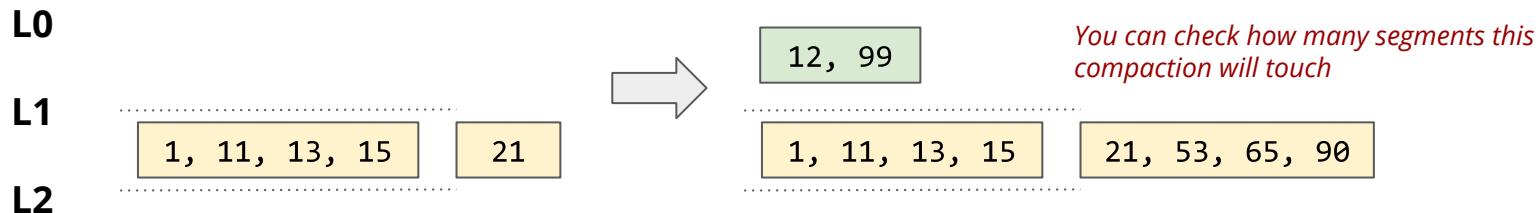


Example Compaction Process

L0 can have duplicates keys in different files



Pick all files which have overlapping ranges



How to Optimize for Searching Files?

Look in: (i) mutable MemTable (ii) look at all the files at L0

- L0 files can contain overlapping key ranges, hence, all files need to be searched at **L0**

Further down, it can be a bit simpler as

- Files at L1 onwards do not have overlapping ranges (they are built that way)
- Hence, for each level, only need to check the range block and the bloom filter, not need to have read the file
- Lower levels contain fresher data (e.g., data at L3 would be newer than at L5)

Also, since indexes are sorted and immutable, support range-based queries

General LSM Considerations

What are the size threshold for each level

What are the block sizes

When to do compaction

- Will result in decreasing the number of files
- Which level should be compacted to which next level
- Also: as L0 fills up the speed of writes will be stalled (in the end it will stop completely)

When to do garbage collection

- Deletion of old values which have been deleted
- Typically read the keys from the tree, and insert them back in the system

An Efficient Design and Implementation of LSM-Tree based Key-Value Store on Open-Channel SSD (2014)

**An Efficient Design and Implementation of
LSM-Tree based Key-Value Store on Open-Channel SSD**

Peng Wang Guangyu Sun
Peking University
{wang-peng, gsun}@pku.edu.cn

Song Jiang *
Peking University and
Wayne State University
sjiang@eng.wayne.edu

Jian Ouyang Shiding Lin
Baidu Inc.
{ouyangjian, linshiding}@baidu.com

Chen Zhang
Peking University
chen.ceca@pku.edu.cn

Jason Cong * †
Peking University and
University of California, Los Angeles
cong@cs.ucla.edu

Abstract
Various key-value (KV) stores are widely employed for data management to support Internet services as they offer higher efficiency, scalability, and availability than relational database systems. The log-structured merge tree (LSM-tree) based KV stores have attracted growing attention because they can eliminate random writes and maintain acceptable read performance. Recently, as the price per unit capacity of NAND flash decreases, solid state disks (SSDs) have been extensively adopted in enterprise-scale data centers to provide high I/O bandwidth and low access latency. However, it is inefficient to naively combine LSM-tree-based KV stores with SSDs, as the high parallelism enabled within the SSD cannot be fully exploited. Current LSM-tree-based KV stores are designed without assuming SSD's multi-channel architecture.

To address this inadequacy, we propose LOCS, a system equipped with a customized SSD design, which exposes its internal flash channels to applications, to work with the LSM-tree-based KV store, specifically LevelDB in this work. We extend LevelDB to explicitly leverage the multi-

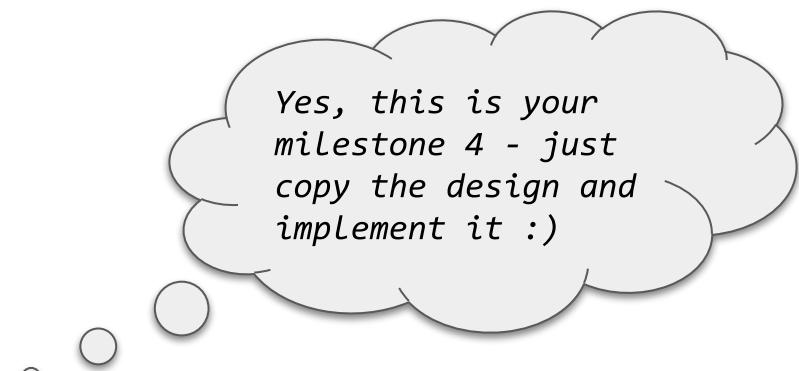
addition, we optimize scheduling and dispatching policies for concurrent I/O requests to further improve the efficiency of data access. Compared with the scenario where a stock LevelDB runs on a conventional SSD, the throughput of storage system can be improved by more than 4× after applying all proposed optimization techniques.

Categories and Subject Descriptors H.3.4 [Information Storage And Retrieval]: Systems and Software

Keywords Solid state disk, flash, key-value store, log-structured merge tree

1. Introduction

With the rapid development of Web 2.0 applications and cloud computing, large-scale distributed storage systems are widely deployed to support Internet-wide services. To store the ultra-large-scale data and service high-concurrent access, the use of traditional relational database management systems (RDBMS) as data storage may not be an efficient choice [15]. A number of features and functionalities of RDBMS, such as transaction consistency guarantee and sup-



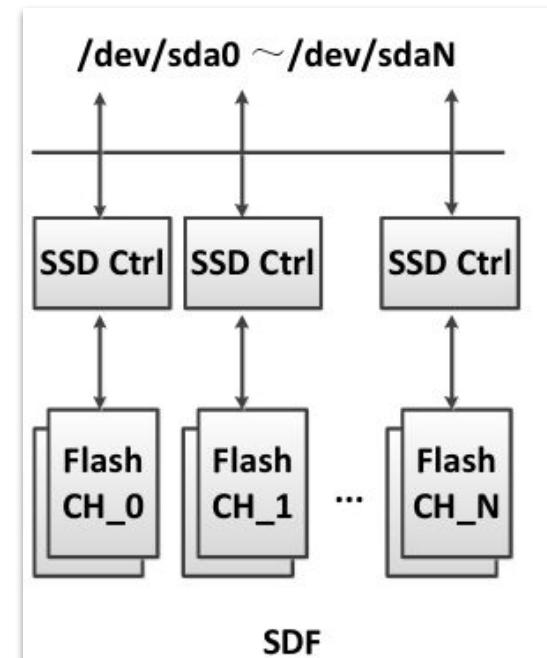
Challenges with the Basic LSM Design

Recall: Open-Channel SSD is similar to SDF where all device internals and placement information is exposed -- high parallelism

1. Single head writing of immutable SSTable
2. Operation unaware scheduling (read, write, erase)
3. Placement and parallelism unaware scheduling

This work: **LOCS**

“LSM-tree-based KV store on Open-Channel SSD”



They retain the basic LSM design, but optimize it for OCSSD

Idea 1: Enable Concurrent Accesses

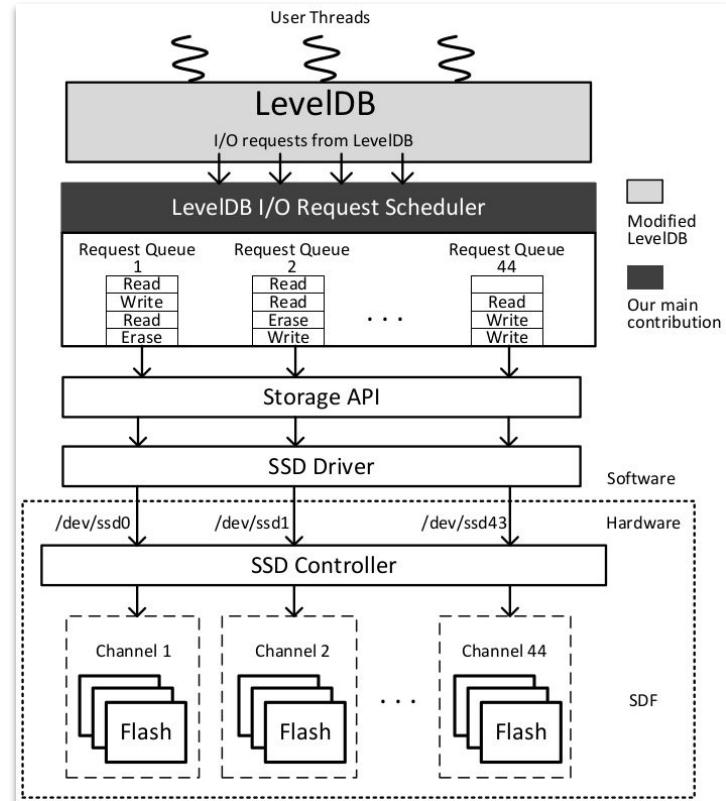
There is still a single mutable MemTable

Number of immutable in-memory MemTables are increased to 44

- Can absorb write bursts

Run multiple parallel compaction at the same time

- Was not possible with HDD because there is only single read/write head



Idea 2: Scheduling Optimization

Question: How should you pick which channel an SSTable should be flushed?

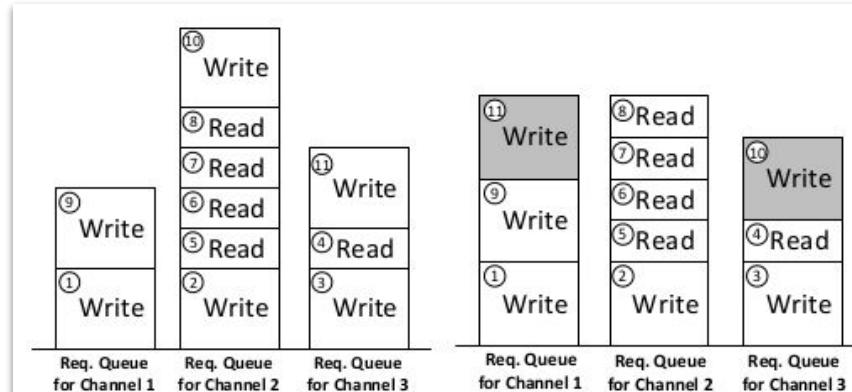
- Writes decides read workload too

Strategy 1: Round-Robin

Strategy 2: Least Weighted Queue Length
Write dispatching

- Weight is read/write/erase cost

$$Length_{weight} = \sum_1^N W_i \times Size_i$$



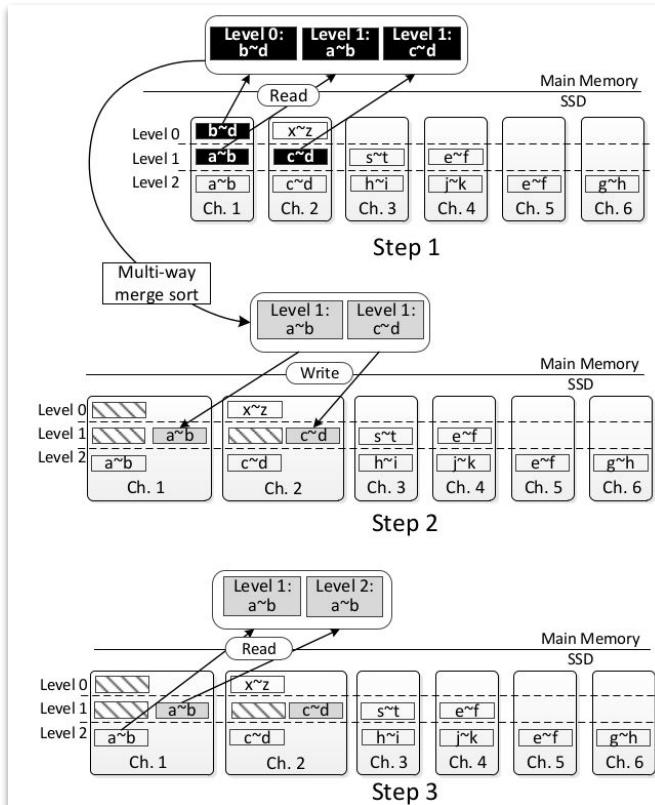
(a) Round-Robin

(b) Least Weighted-Queue-Length

Trace #	1	2	3	4	5	6	7	8	9	10	11
Op.	W	W	W	R	R	R	R	R	W	W	W
Channel	-	-	-	3	2	2	2	2	-	-	-

(c) Trace of the example

Idea 3: Placement Aware Compaction



Recall that LSM trees need compaction

Here: L0 file (b-d) is being pushed to L1

At L1 it overlaps with two files (a-b),(c-d)

We first read those two files in DRAM

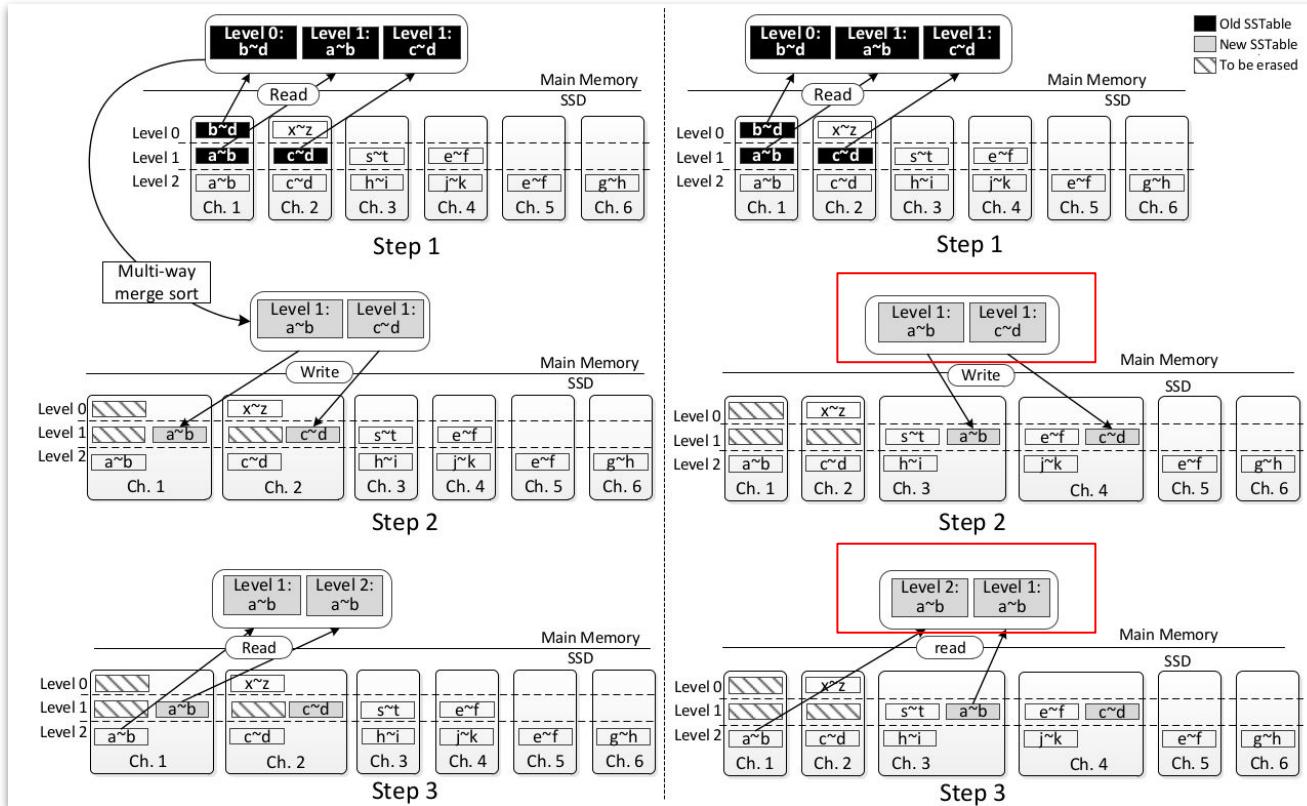
Do a multi-way merge sort with the three files

Then write out the L1 files (a-b) and (c-d)

Next-level of compaction at level L1 and L2 for key ranges of (a-b)

Problem?

Idea 3: Placement Aware Compaction



Idea 4: Erase Aware Scheduling

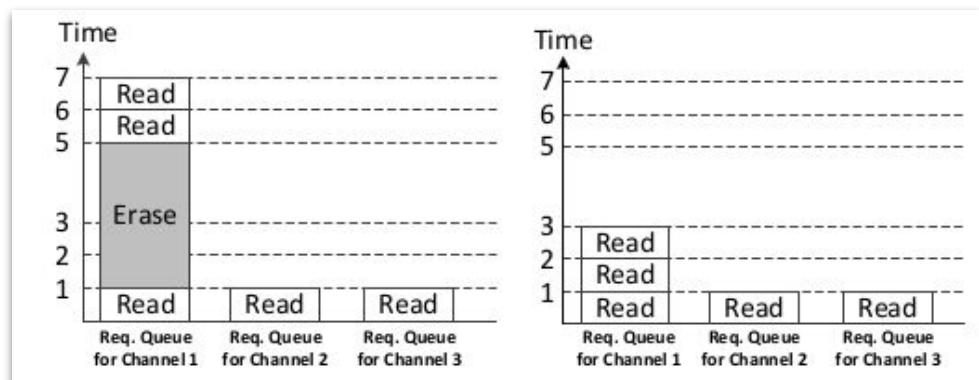
Once the compaction is done, then one must erase blocks

Unlike read/write, erase can be scheduled by the KV when it is most opportune, when is that?

- Eager, as soon as possible

Erase is a long operation

Can lead to interferences with read operation (poor perf)



Eager scheduling of erase might be bad for read performance

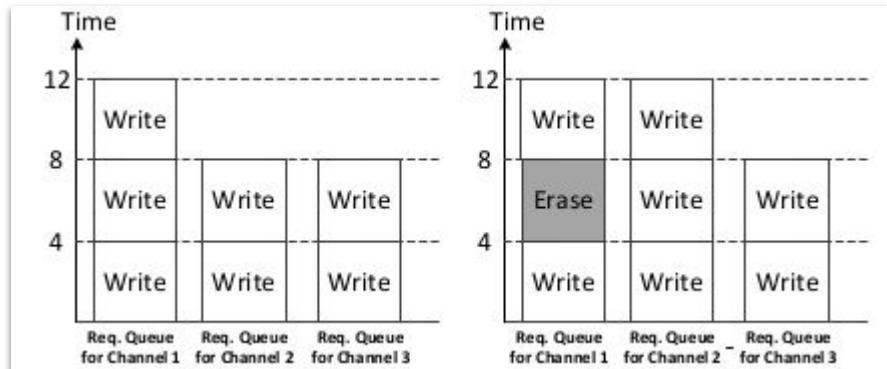
Idea 4: Erase Aware Scheduling

The trick here is to schedule Erase with Writes, not with Read, why?

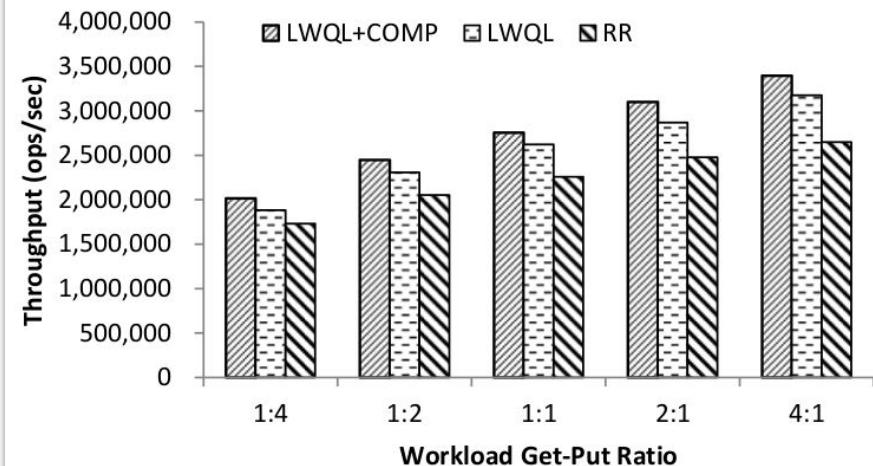
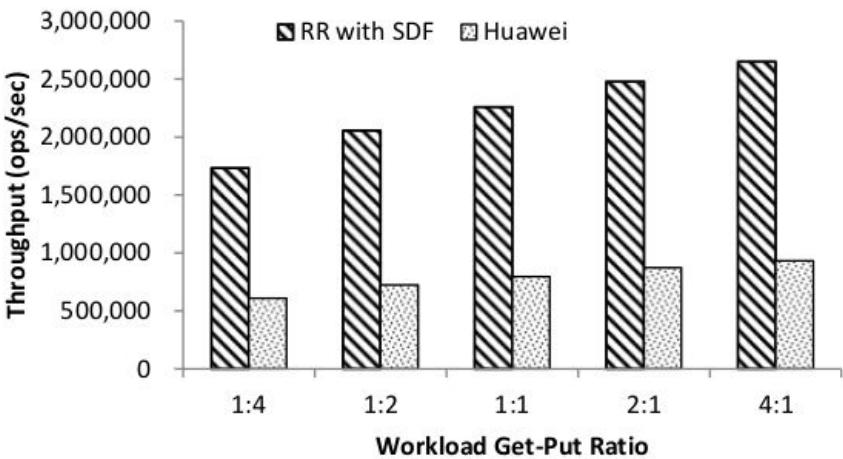
- Because writes can be put to any channel (flexible)
 - Reads cannot be moved around because they need to a read a given address from that channel
- Erase + Write can be used to balance out work among channels

In this example, we can insert Erase with write operations to maintain A balanced LWQL queue

E.g., with Erase in write it will take 19 units, where as Erase in write takes 15 units



Performance: LOCS



Basic idea of software-managed parallelism over channels make sense

RR delivers good performance, LWQL even better, LWQL with Compaction aware optimizations the best of the three

WiscKey: Separating Keys from Values in SSD-Conscious Storage (2016)

WiscKey: Separating Keys from Values in SSD-Conscious Storage

Lanyue Lu, Thanumalayan Sankaranarayana Pillai,
Andrea C. Arpaci-Dusseau, Remzi H. Arpaci-Dusseau

University of Wisconsin, Madison

Abstract

We present WiscKey, a persistent LSM-tree-based key-value store with a performance-oriented data layout that separates keys from values to minimize I/O amplification. The design of WiscKey is highly SSD optimized, leveraging both the sequential and random performance characteristics of the device. We demonstrate the advantages of WiscKey with both microbenchmarks and YCSB workloads. Microbenchmark results show that WiscKey is $2.5 \times$ – $111 \times$ faster than LevelDB for loading a database and $1.6 \times$ – $14 \times$ faster for random lookups. WiscKey is faster than both LevelDB and RocksDB in all six YCSB workloads.

1 Introduction

Persistent key-value stores play a critical role in a variety of modern data-intensive applications, including web indexing [16, 48], e-commerce [24], data deduplication [7, 22], photo stores [12], cloud data [32], social networking [9, 25, 51], online gaming [23], messaging [1, 29], software repository [2] and advertising [20]. By enabling efficient insertions, point lookups, and range queries, key-value stores serve as the foundation for this growing group of important applications.

For write-intensive workloads, key-value stores based on Log-Structured Merge-Trees (LSM-trees) [43] have become the state of the art. Various distributed and local stores built on LSM-trees are widely deployed in large-scale production environments, such as BigTable [16] and LevelDB [48] at Google, Cassandra [33], HBase [29] and RocksDB [25] at Facebook, PNUTS [20] at Yahoo!, and Riak [4] at Basho. The main advantage of LSM-

throughout its lifetime; as we show later (§2), this I/O amplification in typical LSM-trees can reach a factor of 50x or higher [39, 54].

The success of LSM-based technology is tied closely to its usage upon classic hard-disk drives (HDDs). In HDDs, random I/Os are over 100× slower than sequential ones [43]; thus, performing additional sequential reads and writes to continually sort keys and enable efficient lookups represents an excellent trade-off.

However, the storage landscape is quickly changing, and modern solid-state storage devices (SSDs) are supplanting HDDs in many important use cases. As compared to HDDs, SSDs are fundamentally different in their performance and reliability characteristics; when considering key-value storage system design, we believe the following three differences are of paramount importance. First, the difference between random and sequential performance is not nearly as large as with HDDs; thus, an LSM-tree that performs a large number of sequential I/Os to reduce later random I/Os may be wasting bandwidth needlessly. Second, SSDs have a large degree of internal parallelism; an LSM built atop an SSD must be carefully designed to harness said parallelism [53]. Third, SSDs can wear out through repeated writes [34, 40]; the high write amplification in LSM-trees can significantly reduce device lifetime. As we will show in the paper (§4), the combination of these factors greatly impacts LSM-tree performance on SSDs, reducing throughput by 90% and increasing write load by a factor over 10. While replacing an HDD with an SSD underneath an LSM-tree does improve performance, with current LSM-tree technology, the SSD's true potential goes largely unrealized.

So, What is the Problem?

We briefly referenced that reading performance on LSM can be problematic

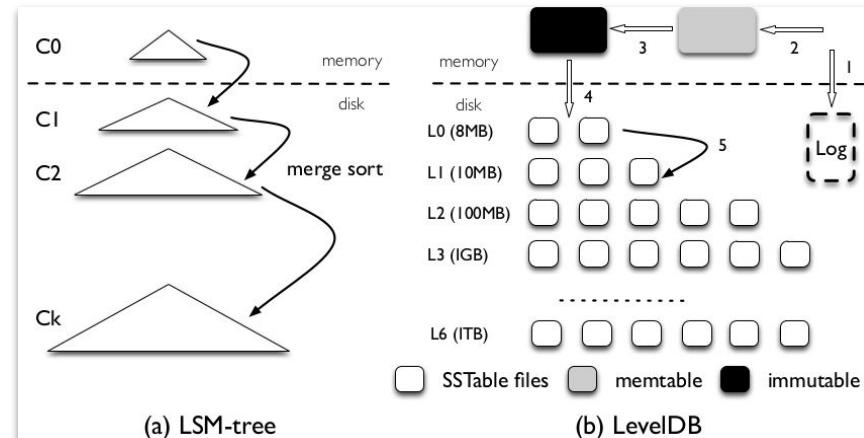
Any guesses why?

What was the read path order?

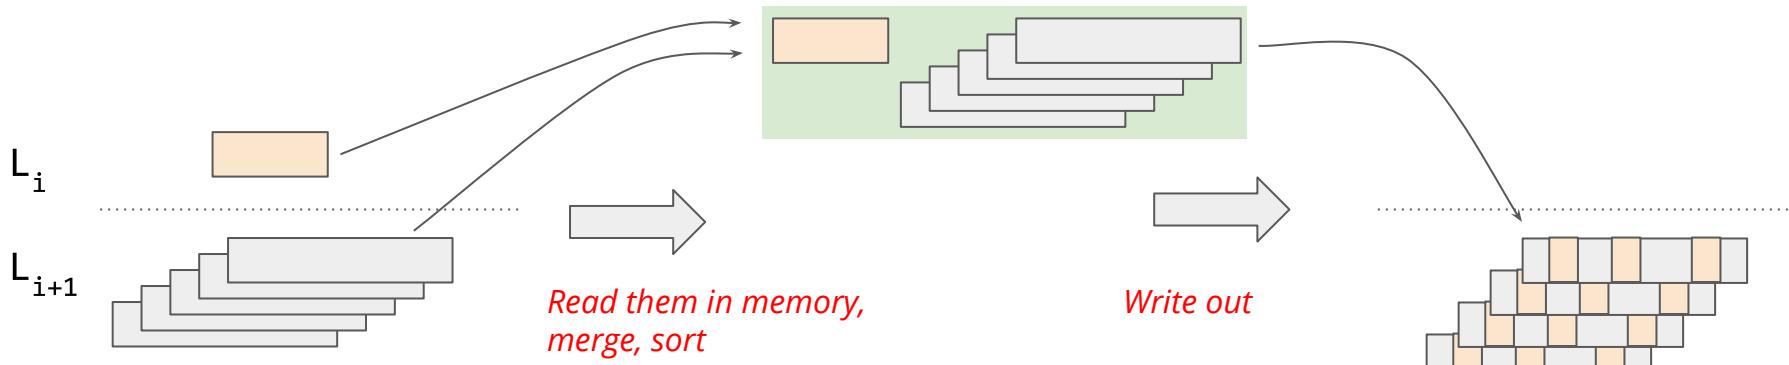
- MemTable → L0 → L1 ... L6 (here)

So, if you were to read simple 1 byte key-value, how much data you have to read before you can find a 1 byte result?

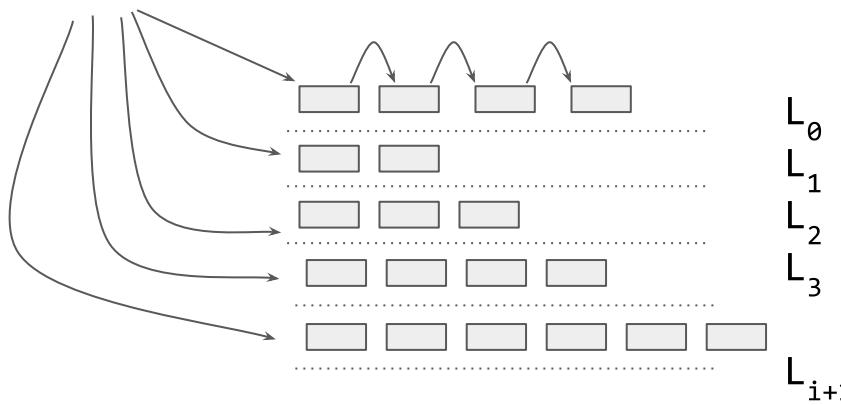
We have looked this type of problem before in FTL for writes
(recall: write-amplification)



LSM has Read and Write Amplifications



read,lookup



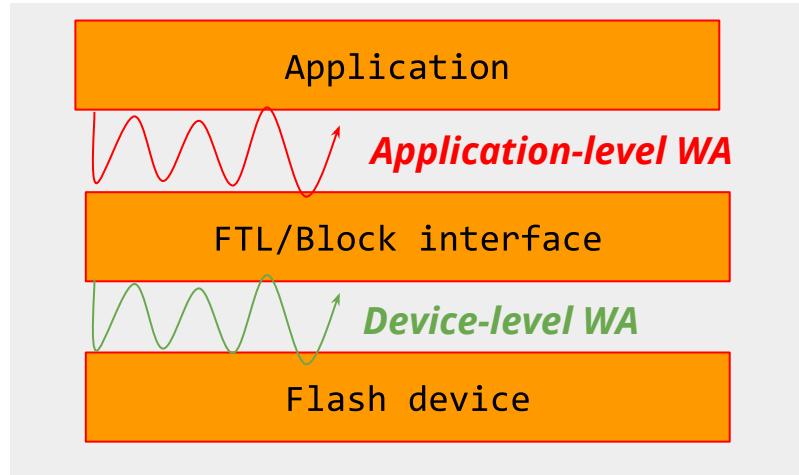
Analysis : Write/Read Amplification (RA/WA)

Compaction can result in

- Reading “n” times data from the next level to merge from the current level
 - For LevelDB this is 10x between levels
 - For 6 levels, it could be 50x

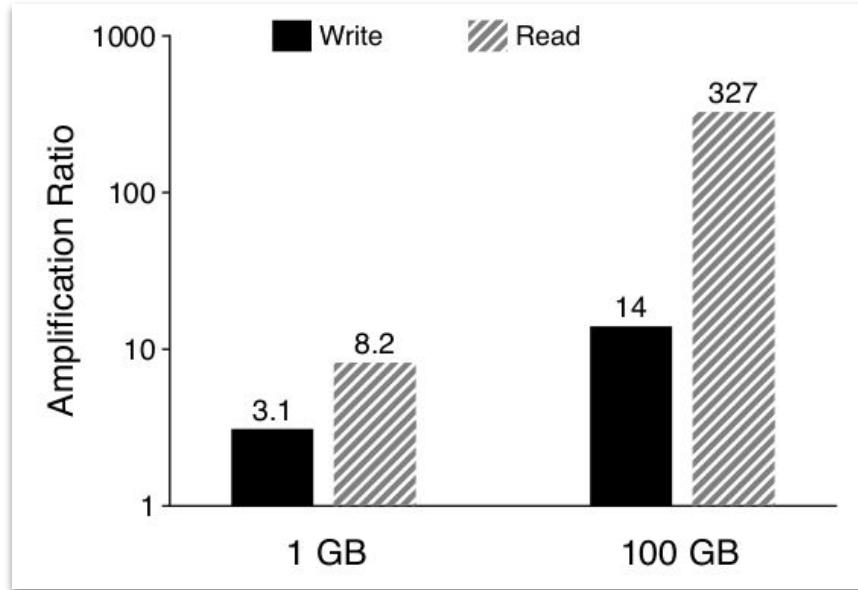
Reading can result in

- Reading “n” files on L0 and then 1 file on following level
 - LevelDB, 8 files (at L0) + 6 files (L1-L6) = 14 files
 - Within the file we need to read the “index” + “bloom filter” + data block
 - For level DB index (16kB), bloom (4kB) + data (4kB)
 - So, if we are looking for a 1kB file: 14 files x (24 kb) = 336 kb \Rightarrow 336x **RA**
 - Determined by how many files do you have to touch and read to find a value



LSM Trees trade high “amplification” for having “sequential performance” \rightarrow Why this was ok with HDDs?

Quantify and Justify



Key size: 16 bytes, value size : 1024 bytes

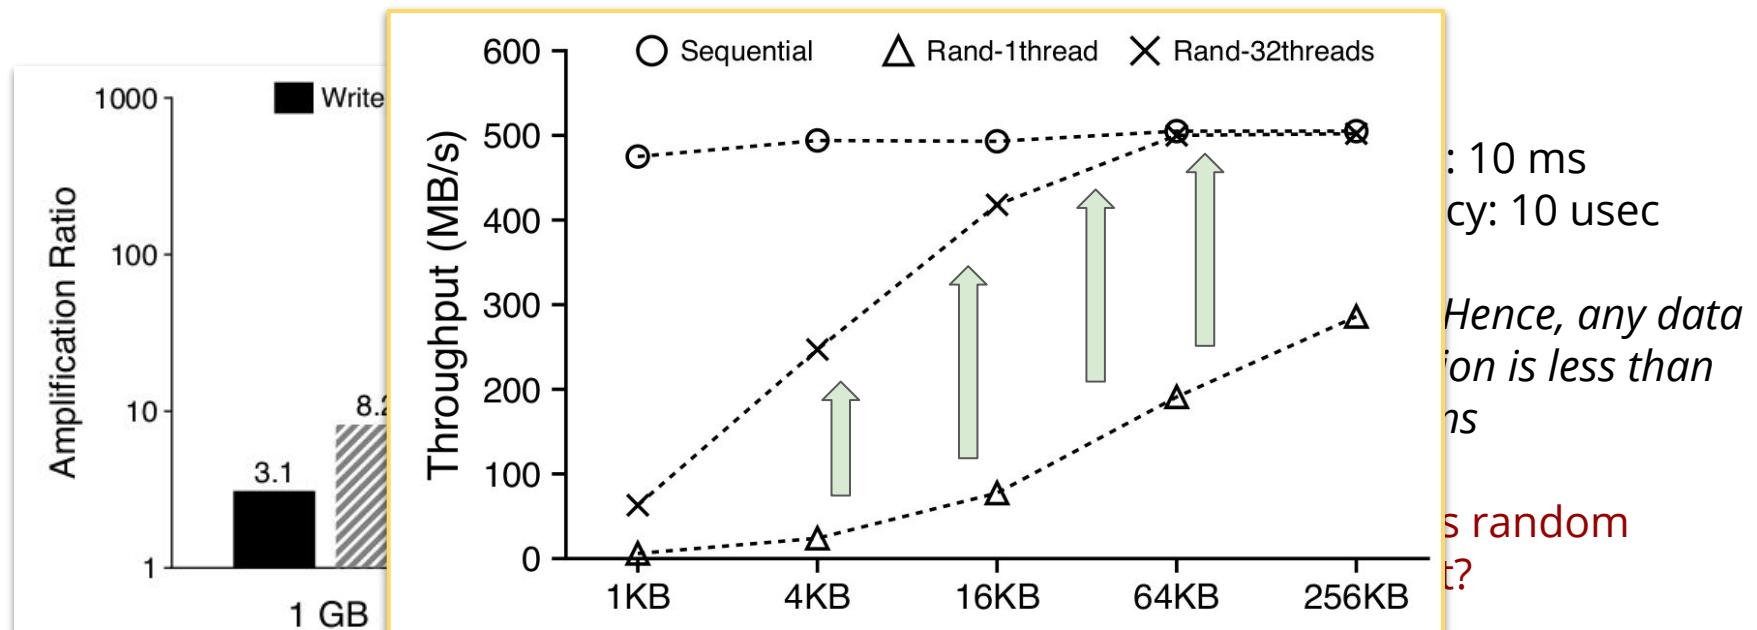
Justification for HDD

- Random 1kB latency: 10 ms
- Sequential 1kB latency: 10 usec

Ratio is seq:rand **1:1000**. Hence, any data structure where amplification is less than 1000, sequential access wins

On SSD? Are sequential vs random accesses are 1:1000 apart?

Quantify and Justify



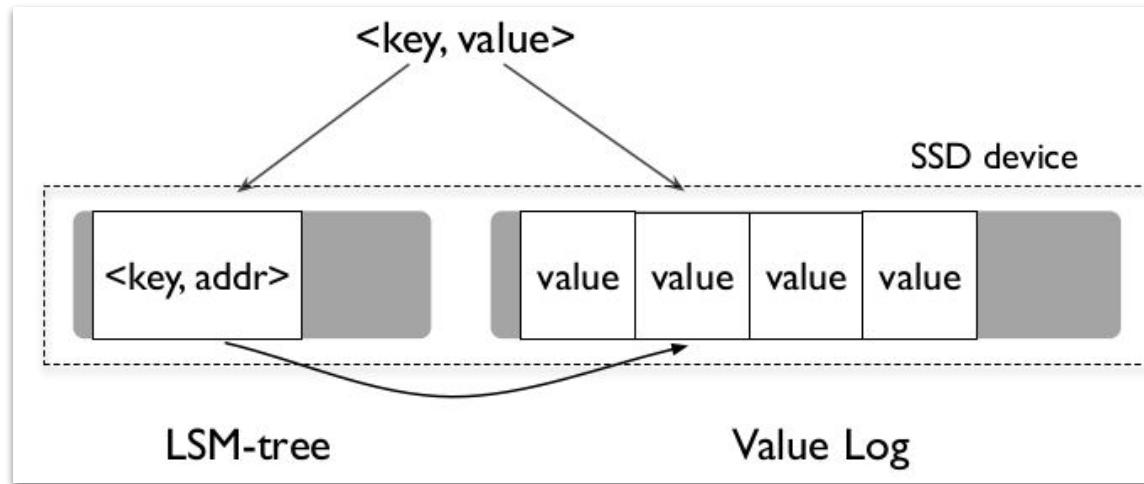
There exists a gap between random and sequential performance, but

- Not for large values
- The gap can be closed by issuing multiple parallel requests

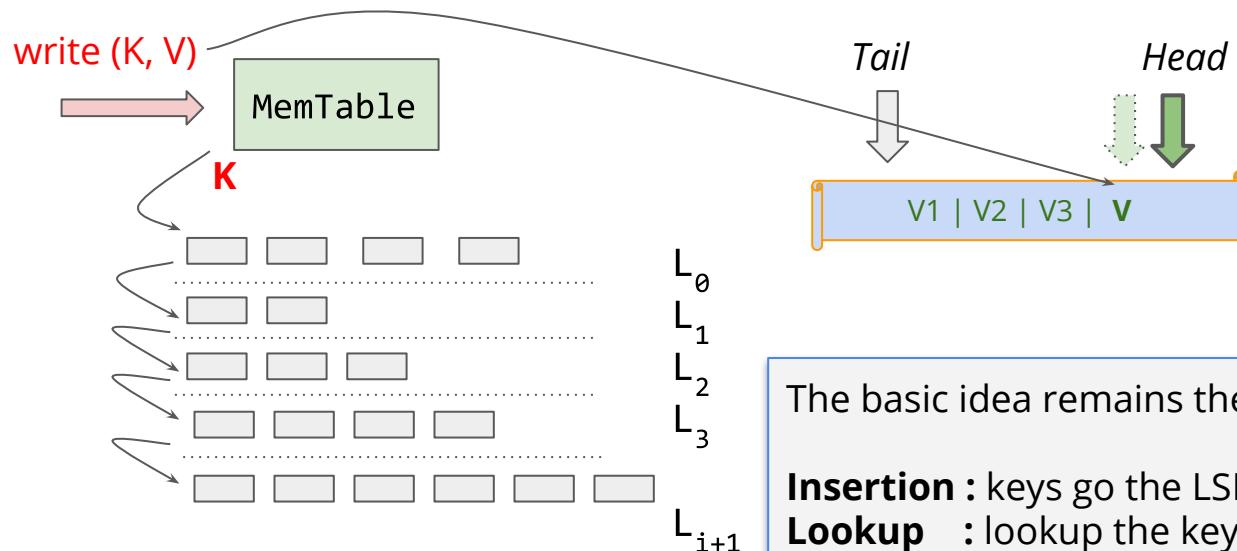
What does WiscKey Proposes

Key Idea: separate keys from the values

- Maintain keys in the LSM tree
- Maintain value in a sequential append value log



Key-Value Insertion and Lookup



The basic idea remains the same

Insertion : keys go the LSM tree, values to the log

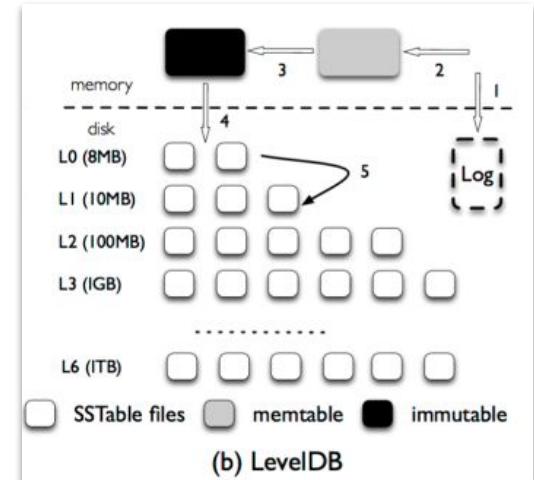
Lookup : lookup the key in the LSM tree, then read the offset from the log

For **range-based queries**, the log can be read in parallel

WiscKey: LSM Tree made out of Keys

What advantages a key-only LSM tree brings

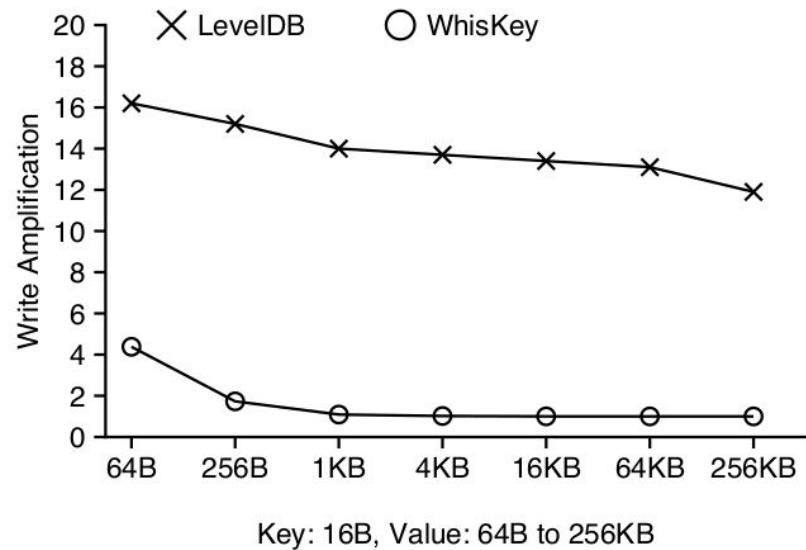
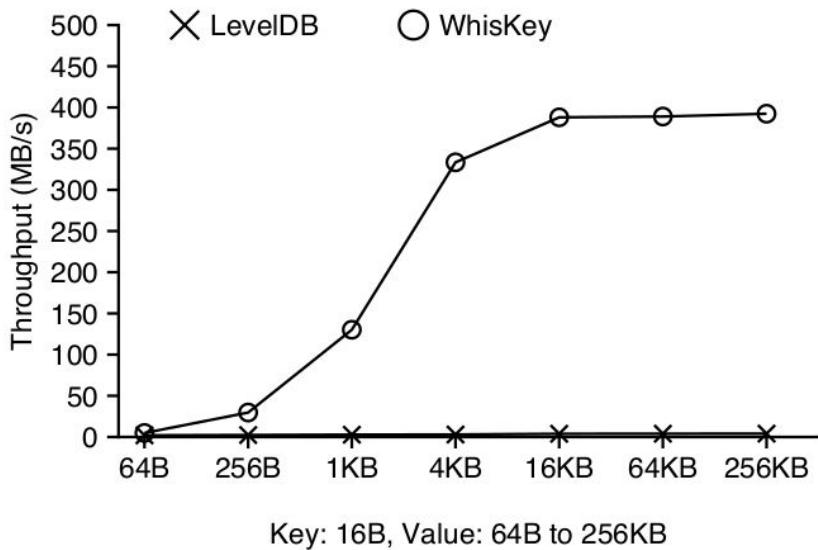
- [with assumptions] keys are small and values are big
- Much improved write-amplification
 - Before WA was: ~10-50x
 - Now $(10 \times \text{key_size}) + \text{value_size} / (\text{key} + \text{value size})$
 - E.g., $(10 \times 16 + 1024) / (1024 + 16) = 1.14$ (not 10x)
 - Worse case : $(50 \times 16 + 1024) / (1024+16) = 1.76$ (not 50x)
- Lower write amplification means longer device life time



Also, the size of the tree can be small (small keys)

- Less levels than a comparable key-value LSM tree
- Small tree can be cached in the memory for fast lookups

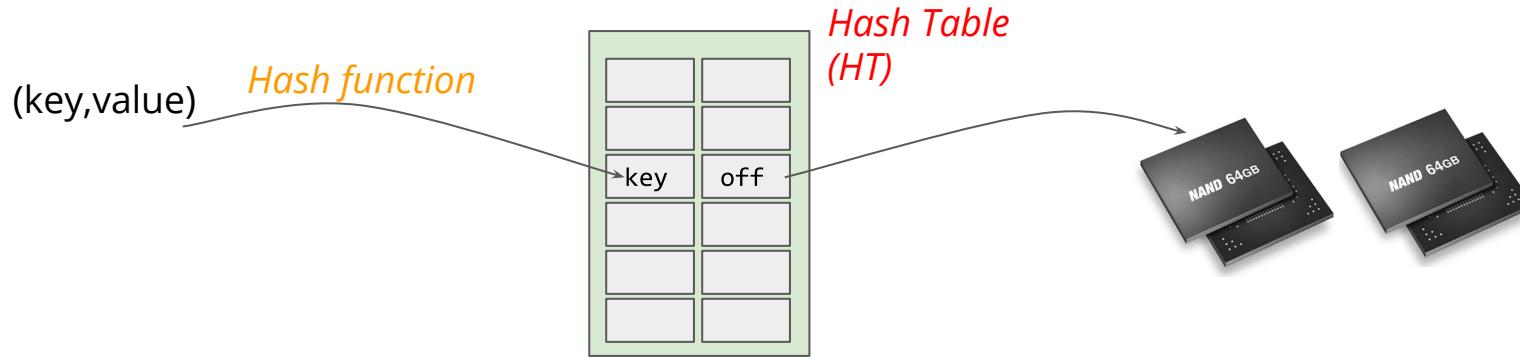
WiscKey: Performance



LevelDB is at 2-4MB/sec whereas WiscKey is at 350 MB/sec (46-111x)

Significantly better write amplification performance

Hash Tables on Flash



This simple hash table based schema works, but it needs to deal with

- Small writes (multiple writes must be packed together)
- Can do fast get and put, but no range-based queries (without additional indexes)
- Trade off {DRAM size of the HT } \longleftrightarrow {number of I/O operations}
 - The same tradeoff as FTL design, how much memory do we need to store a hash table with 1 TB of values
 - Can store the table in flash itself, to decrease the memory size, then multiple I/O

Alternate Hash Table Designs

SkimpyStash: RAM Space Skimpy Key-Value Store on Flash-based Storage

Biplob
1 Mi
EN

ABSTRACT

We present SkimpyStash, a RAM space store on flash-based storage, designed for high-throughput server applications. The distinguishing feature of the design goal is extremely low RAM footprint (0.5) byte per key-value pair, which is more than 10x smaller than prior designs. SkimpyStash uses a hash table to store billions of key-value pairs in a log-structure. To break the barrier of a flash pointer (say, 4 bytes overhead per key), it “moves” most of the pointers from RAM to flash itself. It does this by (i) resolving hash table collisions using linear probing, (ii) storing multiple keys that resolve (collide) to the same slot in a linked list, and (iii) storing the self pointer with a pointer in each hash table bucket. The beginning record of the chain on flash has a flag to indicate if it is a self pointer or not. Two further techniques improve performance: (iv) two-choice based load balancing (to reduce wide variation in bucket sizes), and a bloom filter in each hash table bucket to quickly determine if a key is in RAM to disambiguate the choice during the insertion procedure to pack bucket chain records into flash pages so as to reduce flash reads during bucket size update. The critical design parameter is the bucket size, which is a tradeoff between usage and low lookup latencies. Our evaluation on real-world data center platforms shows that SkimpyStash provides throughputs up to 100,000 get-set operations/sec.

FlashStore: High Throughput Persistent Key-Value Store

Biplob Debnath^{*}
University of Minnesota
Twin Cities, USA
biplob@umn.edu

Sudipta Sengupta
Microsoft Research
Redmond, USA
sudipta@microsoft.com

Jin Li
Microsoft Research
Redmond, USA
jinl@microsoft.com

ABSTRACT

We present FlashStore, a high throughput persistent key-value store, that uses flash memory as a non-volatile cache between RAM and hard disk. FlashStore is designed to store the working set of key-value pairs on flash and use one flash read per key lookup. As the working set changes over time, space is made for the current working set by destroying recently unused key-value pairs to hard disk and recycling pages in the flash store. FlashStore organizes key-value pairs in a log-structure on flash to exploit faster sequential write performance. It uses an in-memory hash table to index them, with hash collisions resolved by a variant of cuckoo hashing. The in-memory hash table stores compact key signatures instead of full keys so as to strike tradeoffs between RAM usage and fast flash read operations.

FlashStore can be used as a high throughput persistent key-value storage layer for a broad range of server class applications. We compare FlashStore with BerkeleyDB, an embedded key-value store application, running on hard disk and flash separately, so as to bring out the performance gain of FlashStore in not only using flash as a cache above hard disk but also in its use of flash aware algorithms. We use real-world data traces from two data center applications, namely, Xbox LIVE Primestream online multi-player game and inline storage deduplication, to drive and evaluate the design of FlashStore on traditional and low power server platforms. FlashStore outperforms BerkeleyDB by up to 60x on throughput (ops/sec), up to 50x on energy efficiency (ops/Joule), and up to 85x on cost efficiency (ops/sec/dollar) on the evaluated datasets.

SILT: A Memory-Efficient, High-Performance Key-Value Store

Hyeontaek Lim¹, Bin Fan¹, David G. Andersen¹, Michael Kaminsky²

¹Carnegie Mellon University, ²Intel Labs

ABSTRACT

SILT (Small Index Large Table) is a memory-efficient, high-performance key-value store system based on flash storage that scales to serve billions of key-value items on a single node. It requires only 0.7 bytes of DRAM per entry and retrieves key/value pairs using on average 1.01 flash reads each. SILT combines new algorithmic and systems techniques to balance the use of memory, storage, and computation. Our contributions include: (1) the design of three basic key-value stores each with a different emphasis on memory-efficiency and write-friendliness; (2) synthesis of the basic key-value stores to build a SILT key-value store system; and (3) an analytical model for tuning system parameters carefully to meet the needs of different workloads. SILT requires one to two orders of magnitude less memory to provide comparable throughput to current high-performance key-value systems on a commodity desktop system with flash storage.

Categories and Subject Descriptors

D.4.2 [Operating Systems]: Storage Management; D.4.7 [Operating Systems]: Organization and Design; D.4.8 [Operating Systems]: Performance; E.1 [Data]; Data Structures; E.2 [Data]; Data Storage Representations; E.4 [Data]: Coding and Information Theory

General Terms

Algorithms, Design, Measurement, Performance

Keywords

Algorithms, design, flash, measurement, memory efficiency, performance

1. INTRODUCTION

Key-value storage systems have become a critical building block for today’s large-scale, high-performance data-intensive applications.

Metric	2008 → 2011	Increase
CPU transistors	731 → 1,170 M	60 %
DRAM capacity	0.062 → 0.153 GB/S	147 %
Flash capacity	0.134 → 0.428 GB/S	219 %
Disk capacity	4.92 → 15.1 GB/S	207 %

Table 1: From 2008 to 2011, flash and hard disk capacity increased much faster than either CPU transistor count or DRAM capacity.

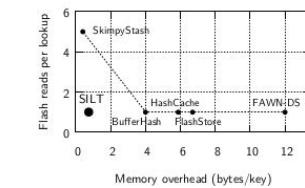


Figure 1: The memory overhead and lookup performance of SILT and the recent key-value stores. For both axes, smaller is better.

e-commerce platforms [21], data deduplication [1, 19, 20], picture stores [7], web object caching [4, 30], and more.

To achieve low latency and high performance, and make best use of limited I/O resources, key-value storage systems require efficient indexes to locate data. As one example, Facebook engineers recently created a new key-value storage system that makes aggressive use of DRAM-based indexes to avoid the bottleneck caused by multiple disk operations when reading data [17]. Unfortunately, DRAM is up to 8X more expensive and uses 25X more power per bit than flash, and as Table 1 shows, is growing more slowly than the capacity of

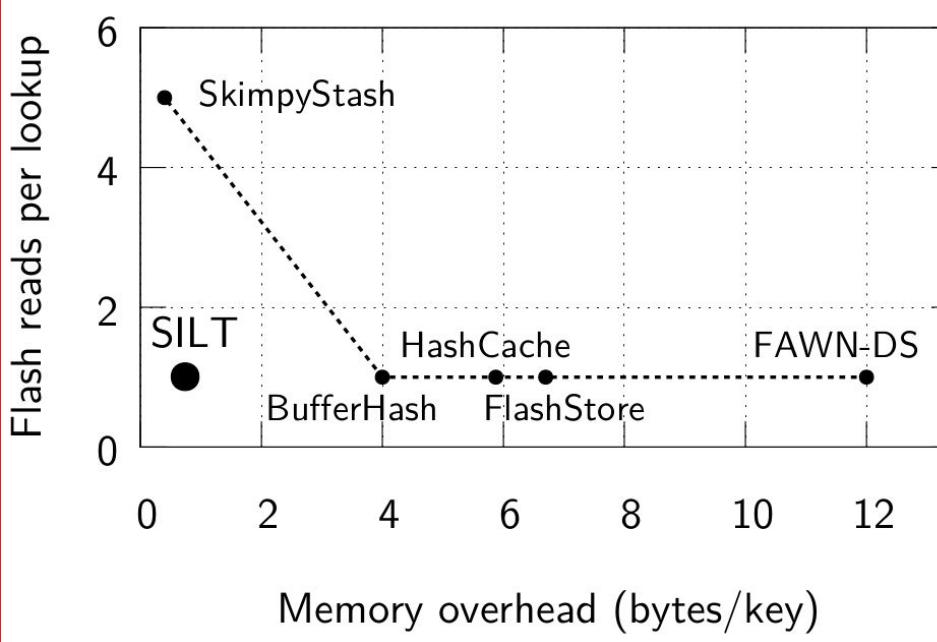
Alternate Hash Table Designs

SkimpyStash: RAM Space Skimpy Key-Value Store on Flash-based Storage

Biplab
Mihai
EN

ABSTRACT

We present SkimpyStash, a RAM space store on flash-based storage, designed for high-throughput server applications. The distinguishing feature of the design goal is extremely low RAM footprint (0.5 byte per key-value pair), which is more efficient than previous designs. SkimpyStash uses a hash table to index key-value pairs stored in a log-structured file. To break the barrier of a flash pointer (say, 4 bytes overhead per key), it “moves” most of the pointer from RAM to flash itself. It does this by (i) resolving hash table collisions using linear probing, (ii) storing multiple keys that resolve (collide) to the same slot in a linked list, and (iii) storing the self pointer with a pointer in each hash table bucket. At the beginning of the chain on flash, there are 6 flash reads per lookup. Two further techniques improve performance: (i) two-choice based load balancing (which reduces wide variation in bucket sizes, hence, chain length and lookup times), and a bloom filter in each hash table entry in RAM to disambiguate the choice during insertion procedure to pack bucket chain records into flash pages so as to reduce flash reads during insertions. Bucket size is the critical design parameter and a powerful knob for making a continuum of tradeoffs between memory usage and low lookup latencies. Our evaluation shows that SkimpyStash provides throughputs up to 100,000 get-set operations/sec.



SILT: A Memory-Efficient, High-Performance Key-Value Store

Bin Fan¹, David G. Andersen¹, Michael Kaminsky²

Carnegie Mellon University, ²Intel Labs

Metric	2008 → 2011	Increase
CPU transistors	731 → 1,170 M	60 %
DRAM capacity	0.062 → 0.153 GB/S	147 %
Flash capacity	0.134 → 0.428 GB/S	219 %
Disk capacity	4.92 → 15.1 GB/S	207 %

Table 1: From 2008 to 2011, flash and hard disk capacity increased much faster than either CPU transistor count or DRAM capacity.

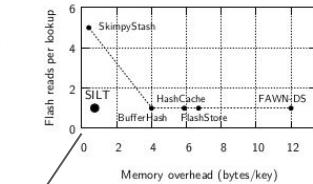


Figure 1: The memory overhead and lookup performance of SILT and the recent key-value stores. For both axes, smaller is better.

e-commerce platforms [21], data deduplication [1, 19, 20], picture stores [7], web object caching [4, 30], and more.

To achieve low latency and high performance, and make best use of limited I/O resources, key-value storage systems require efficient indexes to locate data. As one example, Facebook engineers recently created a new key-value storage system that makes aggressive use of DRAM-based indexes to avoid the bottleneck caused by multiple disk operations when reading data [1]. Unfortunately, DRAM is up to 8X more expensive and uses 25X more power per bit than flash, and, as Table 1 shows, is growing more slowly than the capacity of

the design of FlashStore on traditional and low power server platforms. FlashStore outperforms BerkeleyDB by up to 60x on throughput (ops/sec), up to 50x on energy efficiency (ops/Joule), and up to 85x on cost efficiency (ops/sec/dollar) on the evaluated datasets.

FlashStore, a high performance key-value storage system using flash as a cache between RAM and hard disk. When a key-value blob is written, it is sequentially logged in flash. A specialized RAM-space efficient hash table index using a variant of cuckoo hashing [32] and compact key signatures is used to index the key-value blobs stored in flash mem-

Summary of Data Structures

- B+ Tree (read-optimized)
 - Fast, bounded lookup for read/get ($\log(n)$)
 - Efficient range based queries
 - But poor performance for write-heavy workloads, update bubbling (also small updates)
- Log-structured Merge (LSM) Tree (write-optimized)
 - Good performance for write-heavy workloads, large sequential log based updates
 - Ranged based queries possible
 - Read/Write amplification is a problem
- Simple hash table (hash like md5 on the key → map to a location)
 - [Typically uses] Log-based writing
 - Easy and fast lookup and retrieval ($O(1)$)
 - Limited range based query support (need additional indexing)
 - Tradeoff between (memory usage, and flash I/O)

What you should know from this lecture

1. The idea of B+ Tree, LSM Tree, and Hash Tables
2. Choices these data structures (B+ Tree, LSM, and Hash Table)
3. What advantages and disadvantages they offer when implementing them over NAND flash
4. Key problem and solution: uTree
5. Key problem and solution: LOCS
6. Key problem and solution: WiscKey
7. What is read/write amplification in LSM tree

Further References

- Hyeontaek Lim, Bin Fan, David G. Andersen, and Michael Kaminsky. 2011. SILT: a memory-efficient, high-performance key-value store. In Proceedings of the Twenty-Third ACM Symposium on Operating Systems Principles (SOSP '11).
- Assaf Eisenman, Asaf Cidon, Evgenya Pergament, Or Haimovich, Ryan Stutsman, Mohammad Alizadeh, and Sachin Katti. 2019. Flashield: a hybrid key-value cache that controls flash write amplification. In Proceedings of the 16th USENIX Conference on Networked Systems Design and Implementation (NSDI'19). USENIX Association, USA, 65–78.
- Oana Balmau, Diego Didona, Rachid Guerraoui, Willy Zwaenepoel, Huapeng Yuan, Aashray Arora, Karan Gupta, and Pavan Konka. 2017. TRIAD: creating synergies between memory, disk and log in log structured key-value stores. In Proceedings of the 2017 USENIX Conference on Usenix Annual Technical Conference (USENIX ATC '17). USENIX Association, USA, 363–375.
- PinK: High-speed In-storage Key-value Store with Bounded Tails, USENIX ATC 2020.
- Jiacheng Zhang, Youyou Lu, Jiwu Shu, and Xiongjun Qin. 2017. FlashKV: Accelerating KV Performance with Open-Channel SSDs. ACM Trans. Embed. Comput. Syst. 16, 5s, Article 139 (October 2017), 19 pages.
- Lanyue Lu, Thanumalayan Sankaranarayana Pillai, Andrea C. Arpaci-Dusseau, and Remzi H. Arpaci-Dusseau, WiscKey: Separating Keys from Values in SSD-conscious Storage, USENIX FAST 2016.
- Leonardo Marmol, Swaminathan Sundararaman, Nisha Talagala, and Raju Rangaswami. 2015. NVMKV: a scalable, lightweight, FTL-aware key-value store. In Proceedings of the 2015 USENIX Conference on Usenix Annual Technical Conference (USENIX ATC '15). USENIX Association, USA, 207–219.
- Zhaoyan Shen, Feng Chen, Yichen Jia, and Zili Shao. 2018. DIDACache: An Integration of Device and Application for Flash-based Key-value Caching. ACM Trans. Storage 14, 3, Article 26 (November 2018), 32 pages.

Example 2: HashTable on Flash

FlashStore: High Throughput Persistent Key-Value Store

Biplab Debnath^{*}
University of Minnesota
Twin Cities, USA
biplob@umn.edu

Sudipta Sengupta
Microsoft Research
Redmond, USA
sudipta@microsoft.com

Jin Li
Microsoft Research
Redmond, USA
jinl@microsoft.com

ABSTRACT

We present FlashStore, a high throughput persistent key-value store, that uses flash memory as a non-volatile *cache* between RAM and hard disk. FlashStore is designed to store the working set of key-value pairs on flash and use one flash read per key lookup. As the working set changes over time, space is made for the current working set by destaging recently unused key-value pairs to hard disk and recycling pages in the flash store. FlashStore organizes key-value pairs in a log-structure on flash to exploit faster sequential write performance. It uses an in-memory hash table to index them, with hash collisions resolved by a variant of cuckoo hashing. The in-memory hash table stores compact key signatures instead of full keys so as to strike tradeoffs between RAM usage and false flash read operations.

FlashStore can be used as a high throughput persistent key-value storage layer for a broad range of server class applications. We compare FlashStore with BerkeleyDB, an embedded key-value store application, running on hard disk and flash separately, so as to bring out the performance gain of FlashStore in not only using flash as a cache above hard disk but also in its use of flash aware algorithms. We use real-world data traces from two data center applications, namely, Xbox LIVE Primetime online multi-player game and inline storage deduplication, to drive and evaluate the design of FlashStore on traditional and low power server platforms. FlashStore outperforms BerkeleyDB by up to 60x on throughput (ops/sec), up to 50x on energy efficiency (ops/Joule), and up to 85x on cost efficiency (ops/sec/dollar) on the evaluated datasets.

A high throughput persistent key-value store can help to improve the performance of such applications. Flash memory is a natural choice for such a store, providing persistency and 100-1000 times lower access times than hard disk. Compared to DRAM, flash access times are about 100 times higher. Flash stands in the middle between DRAM and disk also in terms of cost – it is 10x cheaper than DRAM, while 20x more expensive than disk – thus, making it an ideal gap filler between DRAM and disk.

There are two types of popular flash devices, NOR and NAND flash. NAND flash architecture allows a denser layout and greater storage capacity per chip. As a result, NAND flash memory has been significantly cheaper than DRAM, with costs decreasing at faster speeds. NAND flash characteristics have led to an explosion in its usage in consumer electronic devices, such as MP3 players, phones, caches and Solid State Disks (SSDs). In the rest of the paper, we use NAND flash based SSDs as the architectural choice and simply refer to it as flash memory. We describe SSDs in detail in Section 2. To get the maximum performance per dollar out of SSDs, it is necessary to use flash aware data structures and algorithms to avoid small random writes that not only have a higher latency but also reduce flash device lifetimes through increased page wearing.

In this paper, we present the design and evaluation of FlashStore, a high performance key-value storage system using flash as a cache between RAM and hard disk. When a key-value blob is written, it is sequentially logged in flash. A specialized RAM-space efficient hash table index using a variant of cuckoo hashing [32] and compact key signatures is used to index the key-value blobs stored in flash mem-

FlashStore: Data Structures

Many workloads are read-heavy and do not need indexing (B+ tree a bit of an overkill) - restrictive layout how the keys can be stored

- Microsoft wanted to have flash SSDs as a KV cache in front of their HDDs

If we just do a simple hash(key) → location, that would be good enough

- Hash has $O(1)$ lookup time, not $O(\log(n))$ like B+ tree

But the “small write” problem. We cannot store each key in its own page (inefficient) and cannot do small writes to just update the key

Goal: fast KV cache with a single flash I/O read to locate data

Design Goals and Issues

1. Deliver low-latency, high-throughput operations
 - a. For small key looks up
 - b. Values can be in DRAM cache or on Flash
2. Use flash-aware data structures
 - a. Do not do small page updates
3. Low RAM footprint for indexing to lookup on flash
 - a. Technically you can use 8 bytes per key and 64 bytes of value
 - b. So for a 1 TB of flash drive, you will need $1 \text{ TB} / (64 + 8) \times 8 \text{ bytes} = 122 \text{ GB}$ of DRAM (!)
 - c. Same problem as with the FTL

Architecture

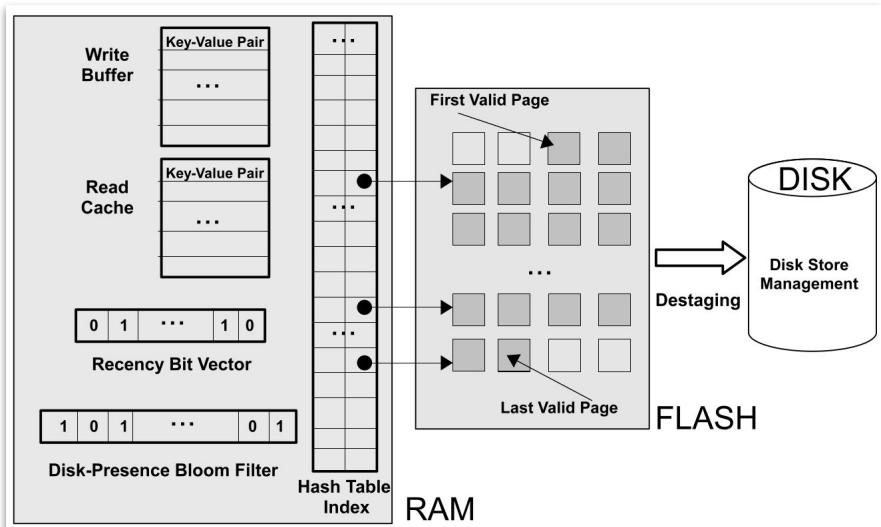
RAM Write buffer : buffer until the flash page size

Read cache: fixed-size read cache for recently used items (LRU)

Recency Bit Vector: maintains access information for staging data between flash and disk

Bloom filter: probabilistic “false positive”, but never “false negative” (*it's not there when it is there*)

HashTable: The primary data structure to look for key → flash location in one flash read



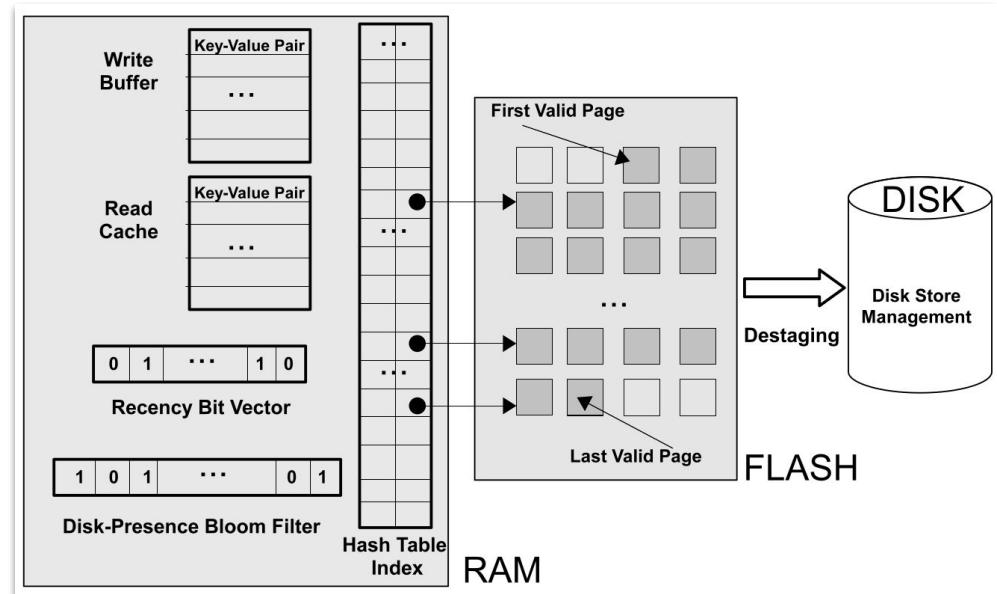
Key Lookup and Insertion Operations

Insert (with timestamps):

1. Into the write buffer
2. Wait until full
3. Write out to flash
4. Update the HT index

Lookup

1. In RAM read cache
2. In RAM write cache
3. Lookup in HT index to find on flash
4. Lookup bloom filter
 - a. No: return NULL
 - b. Yes: disk search (B+ tree)
5. Update recency bit
6. (Optional) put in RAM read cache



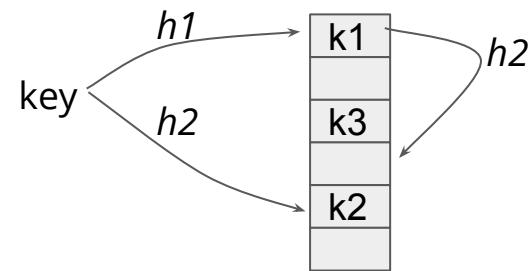
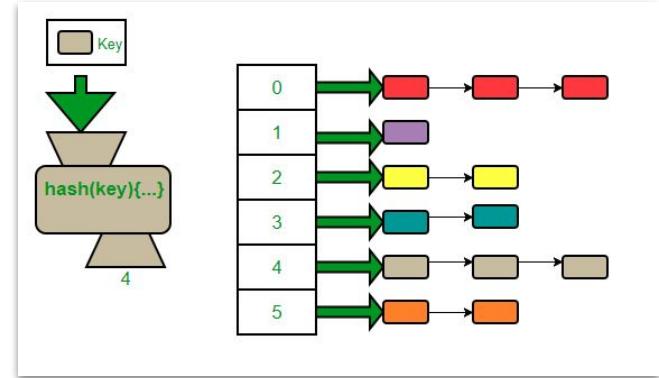
Hash Table Design

In a simple hash table, we can do something like

- Hash(key) → HT slot → check if the key stored there matches
 - OK, then follow the flash page pointer (8bytes)
 - Collision: then follow the link list of collision pointers

Uses **Cuckoo hashing** : use “n” hash functions and find the first free location to put the key. No need to scan any linear list in case of high collision

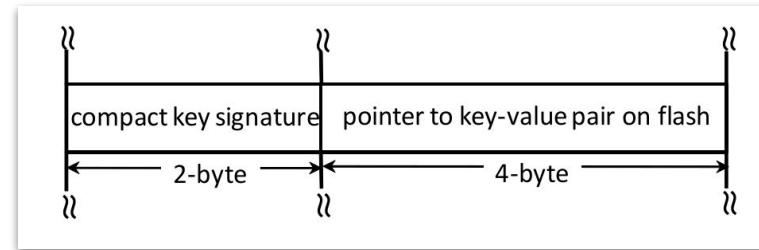
What to store in these hash table slots? Full key and flash page address? (lots of data)



Hash Table Memory Usage: What to Store?

Compact key signature (instead of full key and hash):

- A full key can be of any size, hashes are large too (160-512 bits)
- If the key used i^{th} hash function then used the top-order **16 bits** as a compact signature



Flash page offset as 4 byte pointers (not 8 bytes) : maximum size = $2^{32} \times 4\text{KB} = 8\text{TB}$

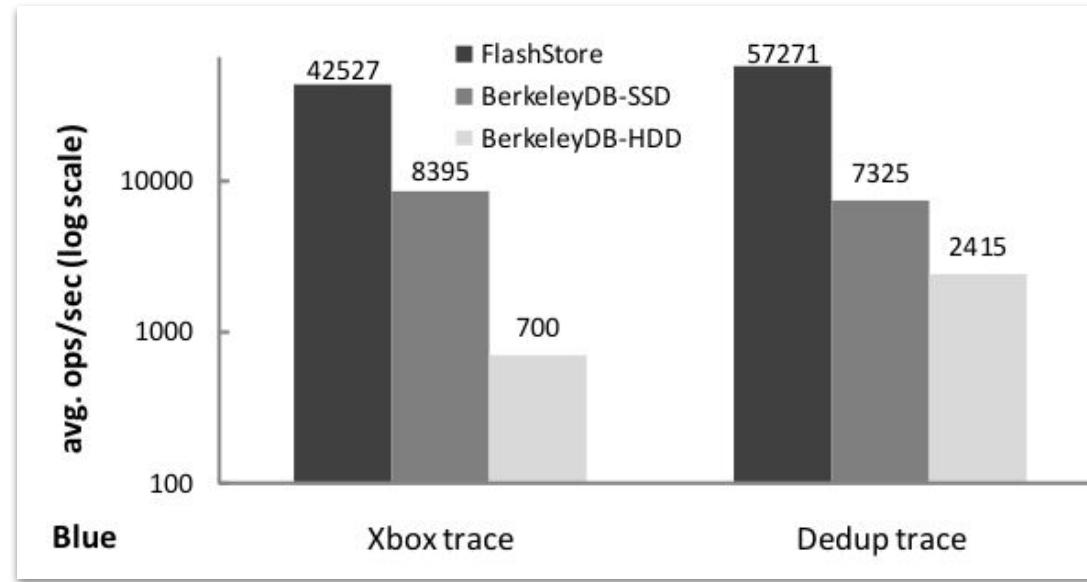
- How many bits to use, can be optimized for the given size of the device
- For example, 160GB device (what they used), $160\text{GB}/4\text{KB} = 26$ bits only
 - Rest of the $(32 - 26) = 6$ bits, can be used for in-page offsets of 128 bytes
 - Hence, 128 bytes becomes the minimum packing granularity

Broadly speaking: a memory-efficient HT table design is an active research problem
(many papers are out there in the field, we are only covering one trick)

Flash Specific Concerns

- Filled flash pages are written in a log-append order (lookup is done using the in-memory HT table)
 - Log garbage collection for entries that have been overwritten or deleted (similar logic)
- After certain HT table occupancy and Flash usage - trigger destaging from flash to HDD
 - Pick pages and check the recency bitmap in memory to find if they have been accessed recently
 - Yes, put them in write buffer (back in the circulation)
 - No, push them to HDD and make space
- At crash
 - Default option: build HT by scanning flash logs
 - Options 2: checkpointing

Performance

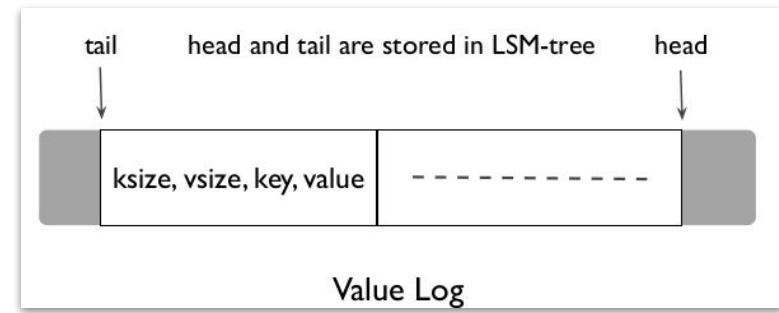


Delivers performance for two important workloads for Microsoft (xbox, and dedup)

Compared with running BerkeleyDB (B+Tree) on SSD and HDD

WiscKey: Doing garbage collection in vLog

A **native way** would be : to scan the LSM key tree to identify all valid values and then remove them.



Better way: to keep a back reference to the keys in the value log as well

Once GC kicks in, values from the tail are read, validated by querying the LSM tree, and then move to the head

The new tail, and addresses are then inserted in the LSM tree before cleaning values